

SAP Certified Associate - Backend Developer - SAP Cloud Application Programming Model

- <https://learning.sap.com/certifications/sap-certified-associate-backend-developer-sap-cloud-programming-model#about>
- Exam: C_CPE_16
- Questions: 80
- Cut Score: 65%

How to prepare ?

Learning Journey • Intermediate • 8 hrs

Developing Full-Stack Applications Using Productivity Tools in SAP Business Application Studio

This Learning Journey introduces you to build and deploy business applications in SAP Business Application Studio using productivity tools.

Learning Journey • Intermediate • 16 hrs

Delivering Side-by-Side Extensibility based on SAP BTP, Kyma runtime

This Learning Journey introduces you into all important aspects to develop and run containerized apps and extensions on SAP BTP, Kyma Runtime

Learning Journey • Intermediate • 8 hrs

Building side-by-side extensions on SAP BTP

This learning journey will guide you through the process of building a side-by-side extension using the SAP Cloud Application Programming Model (Node.js) alongside different development tools an...

Learning Journey • Intermediate • 11 hrs

Developing Advanced Extensions with SAP Cloud SDK

This learning journey allows you to develop side-by-side extensions on SAP Business Technology Platform for SAP solutions like SAP S/4HANA Cloud Public Edition or SAP SuccessFactors with SAP...

How to prepare ?

Learning Journey	Number of hours
Full-Stack applications in SAP Business Application Studio	8 hours
Side-by-Side extensions on SAP BTP	8 hours
Side-by-Side extensibility based on SAP BTP, Kyma runtime	16 hours
Developing advanced extensions with SAP Cloud SDK	11 hours
Total	43 hours

And...

- <https://learning.sap.com/learning-journeys/developing-applications-on-sap-btp-cloud-foundry-runtime>
- Developing Applications on SAP BTP, Cloud Foundry runtime
- 4 hours

Topic Areas

Application Extension Development & Deployment

Exam percentage: 21% - 30%

Application Security

Exam percentage: 11% - 20%

SAP Cloud Application Programming Model

Exam percentage: 11% - 20%

DevOps and Continuous Delivery

Exam percentage: 11% - 20%

SAP Build Process Automation

Exam percentage: ≤10%

SAP S/4HANA Cloud Extensibility

Exam percentage: 11% - 20%

Kyma Runtime

Unit 1 – What is Kubernetes ?

- Open source system for automating deployment, scaling and management of containerized applications
- Can run almost anywhere
 - Public cloud
 - Private cloud
 - On-Premise
- All major cloud providers offer Kubernetes as a managed service
- SAP BTP Kyma runtime is a managed Kubernetes service

Unit 1 – Kubernetes Cluster

Cluster is a set of machines
(physical or virtual)

Each k8s cluster has a master node
and one or more worker nodes

Master node is responsible for
managing the cluster

Worker nodes are responsible for
running the cluster

Master Node (Control Plane)



Worker Node



Worker Node



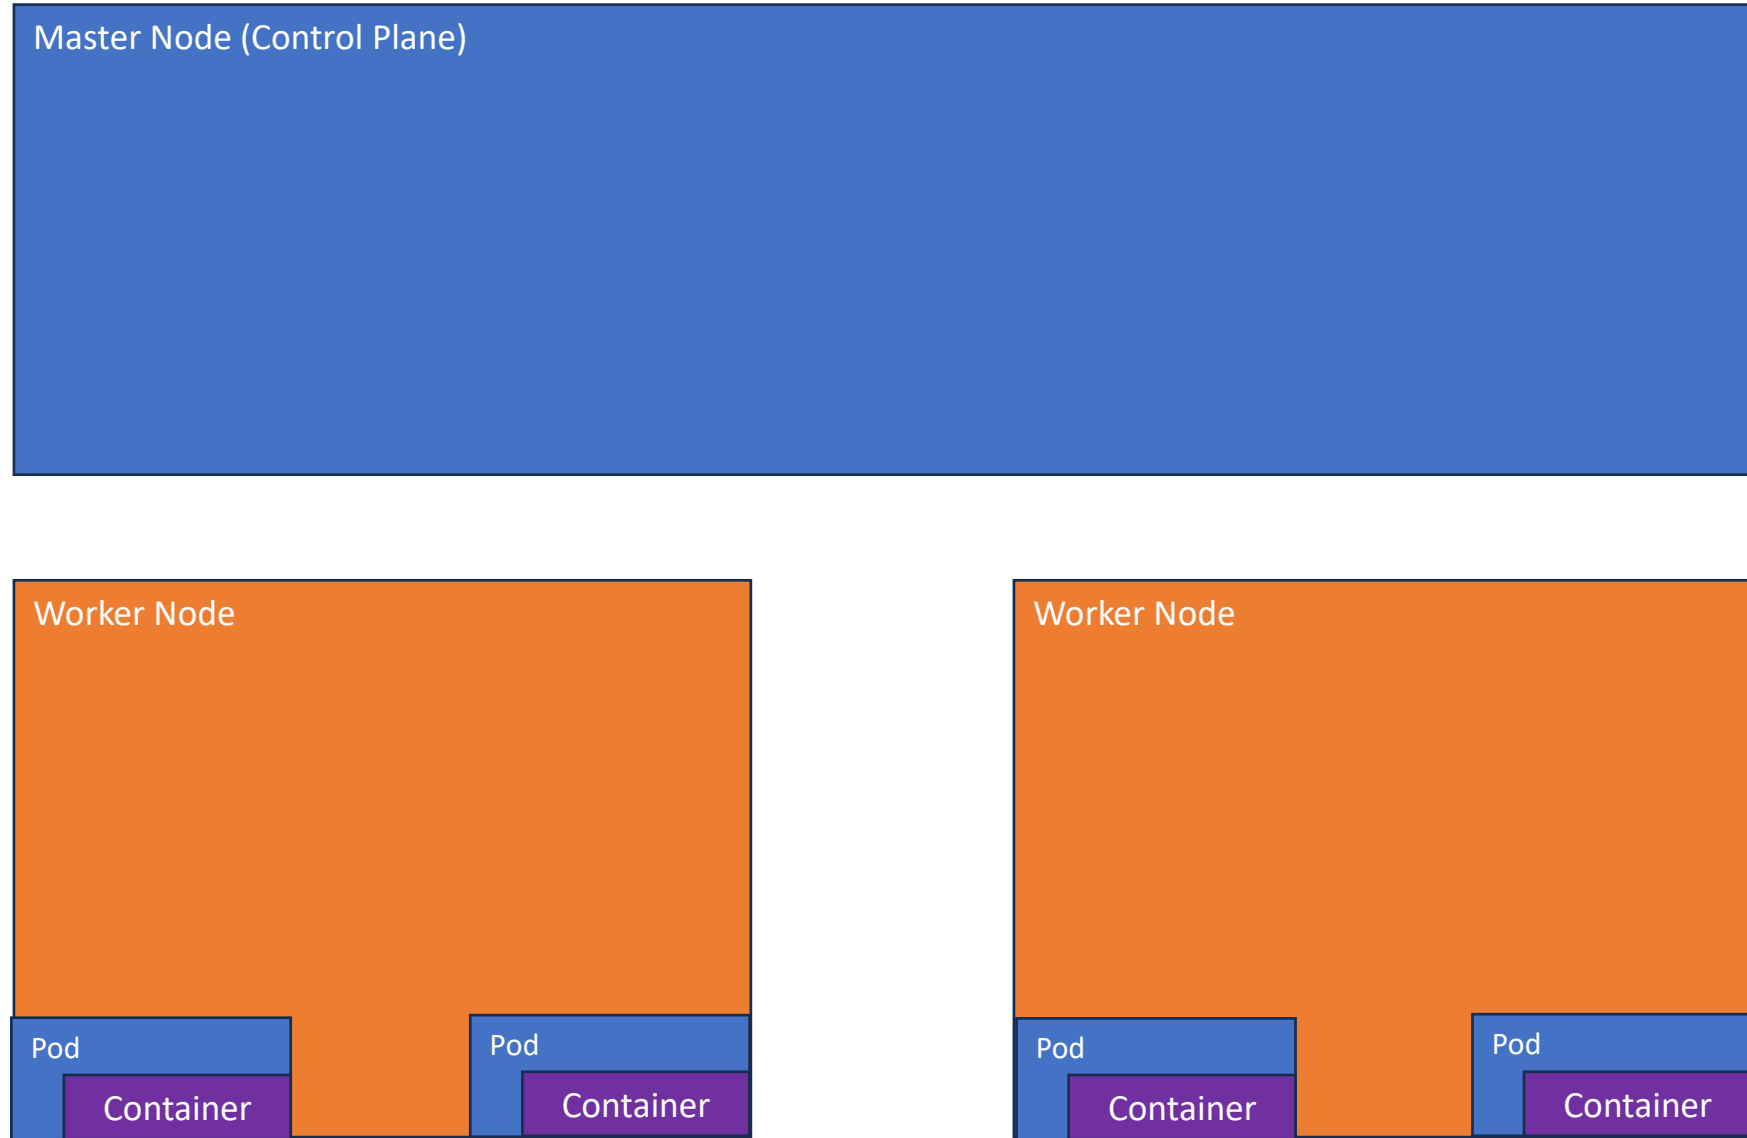
Unit 1 – Kubernetes Cluster

Pods – smallest deployable units in k8s

Thin wrapper for containers

Usually 1 application per Pod

Pods are ephemeral – can die easily



Unit 1 – Kubernetes Cluster

K8s does not want to be tied to a single container technology (for example, Docker)

K8s allows you to use any container runtime that implements this Container Runtime Interface

Master Node (Control Plane)

Worker Node

Container Runtime Interface CRI

Container Runtime

Pod

Container

Pod

Container

Worker Node

Container Runtime Interface CRI

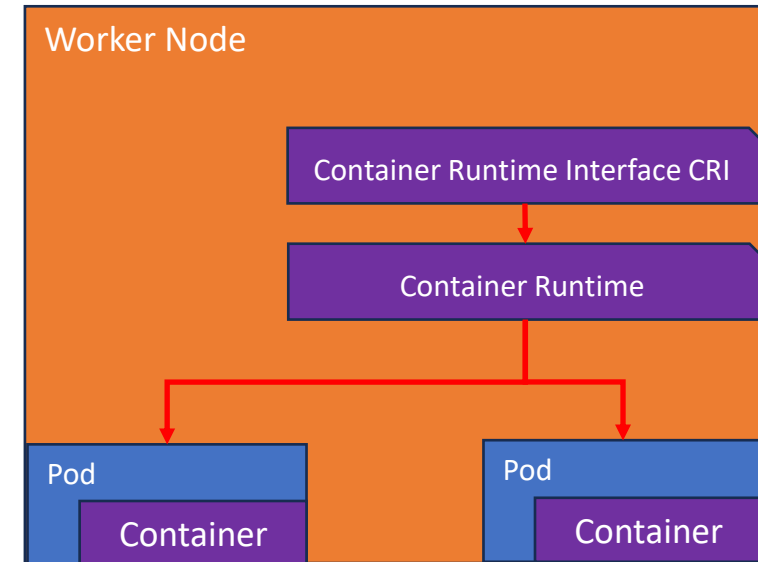

Container Runtime

Pod

Container

Pod

Container



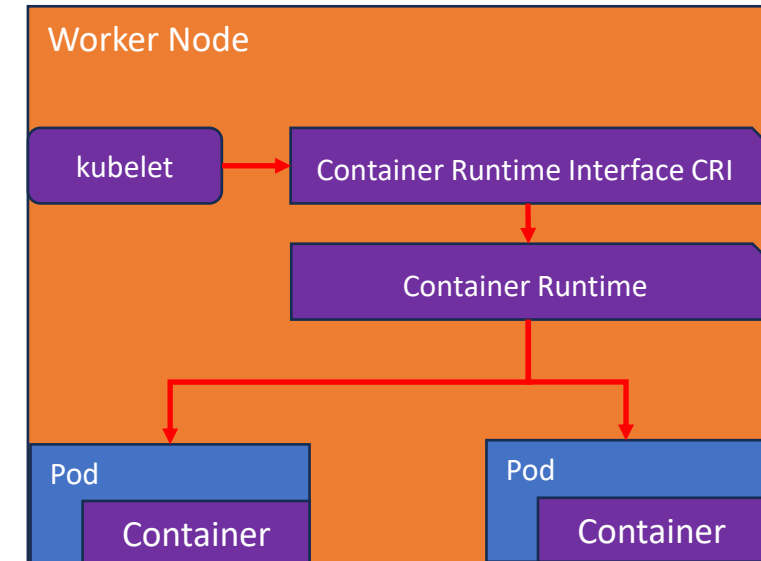
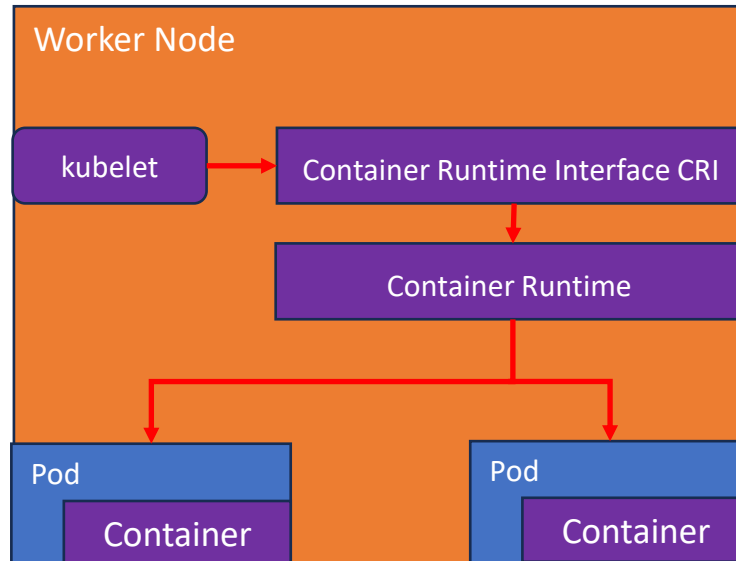
Unit 1 – Kubernetes Cluster

Each worker node has a kubelet process running on it

Ensures containers are running in a pod based on their specs

Communicates with Master Node to receive new work assignments

Master Node (Control Plane)



Unit 1 – Kubernetes Cluster

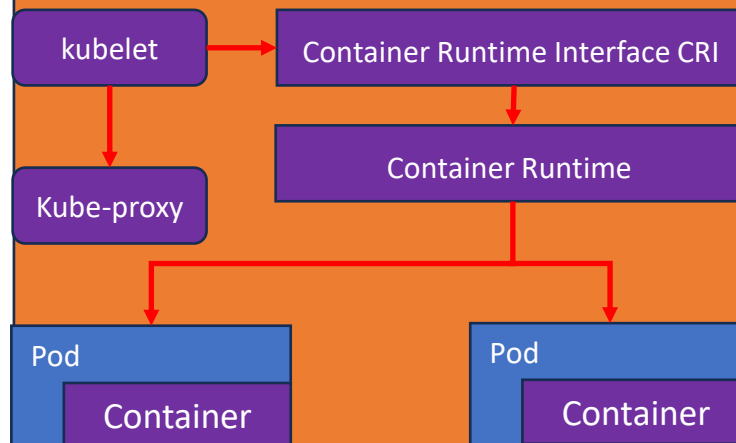
kube-proxy is a network proxy that runs on each node

Main job is to maintain network rules on nodes

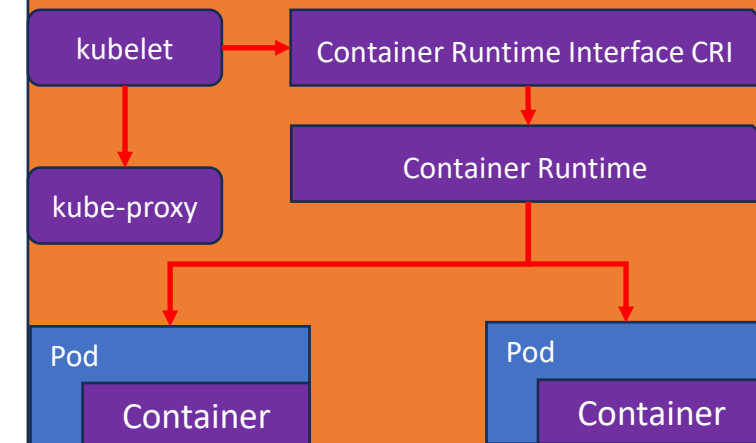
Allows for internal and external cluster communication with pods

Master Node (Control Plane)

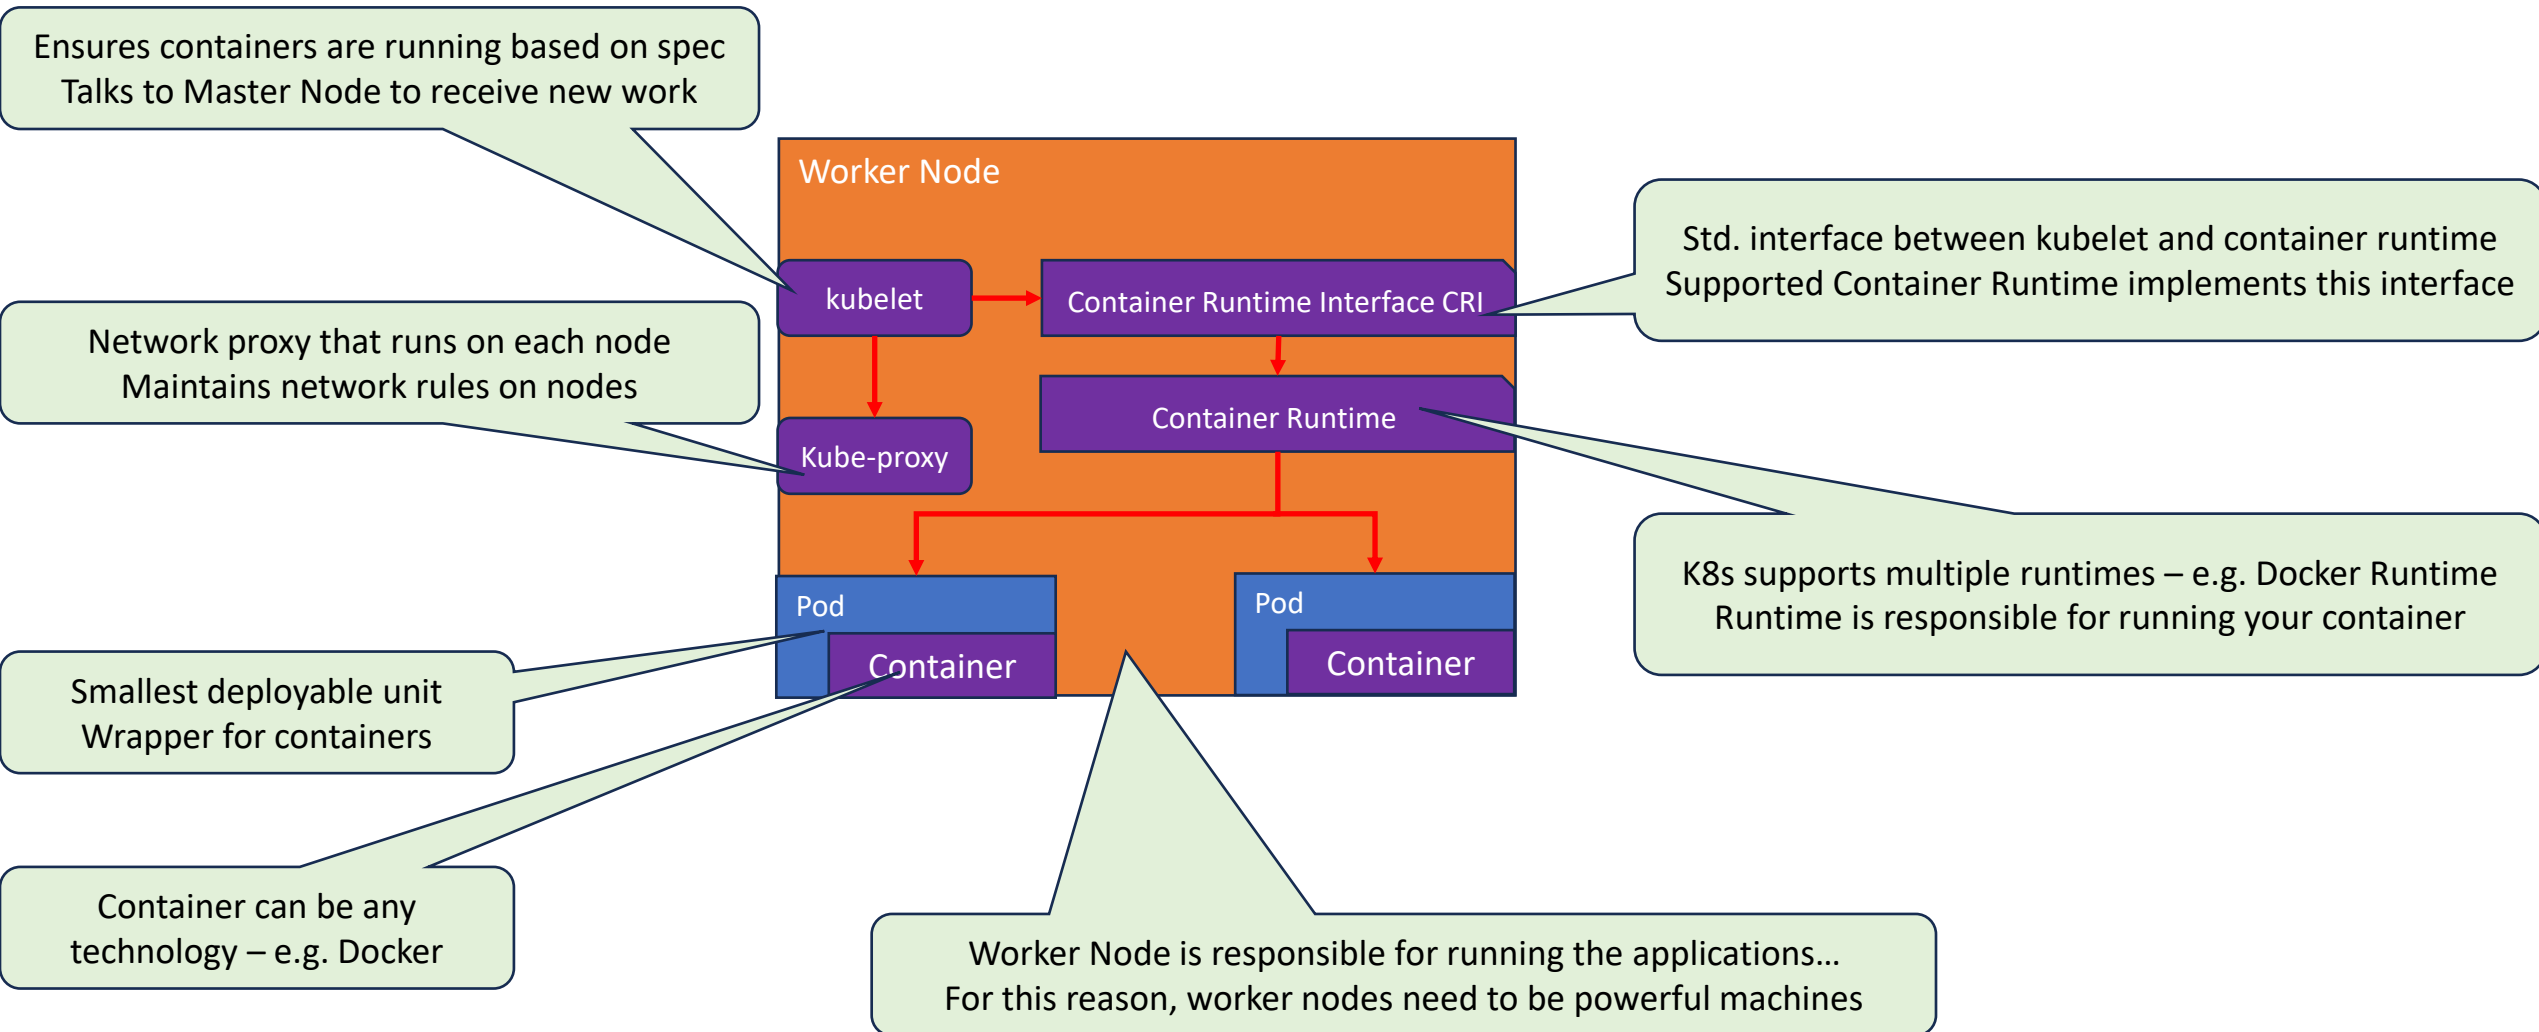
Worker Node



Worker Node



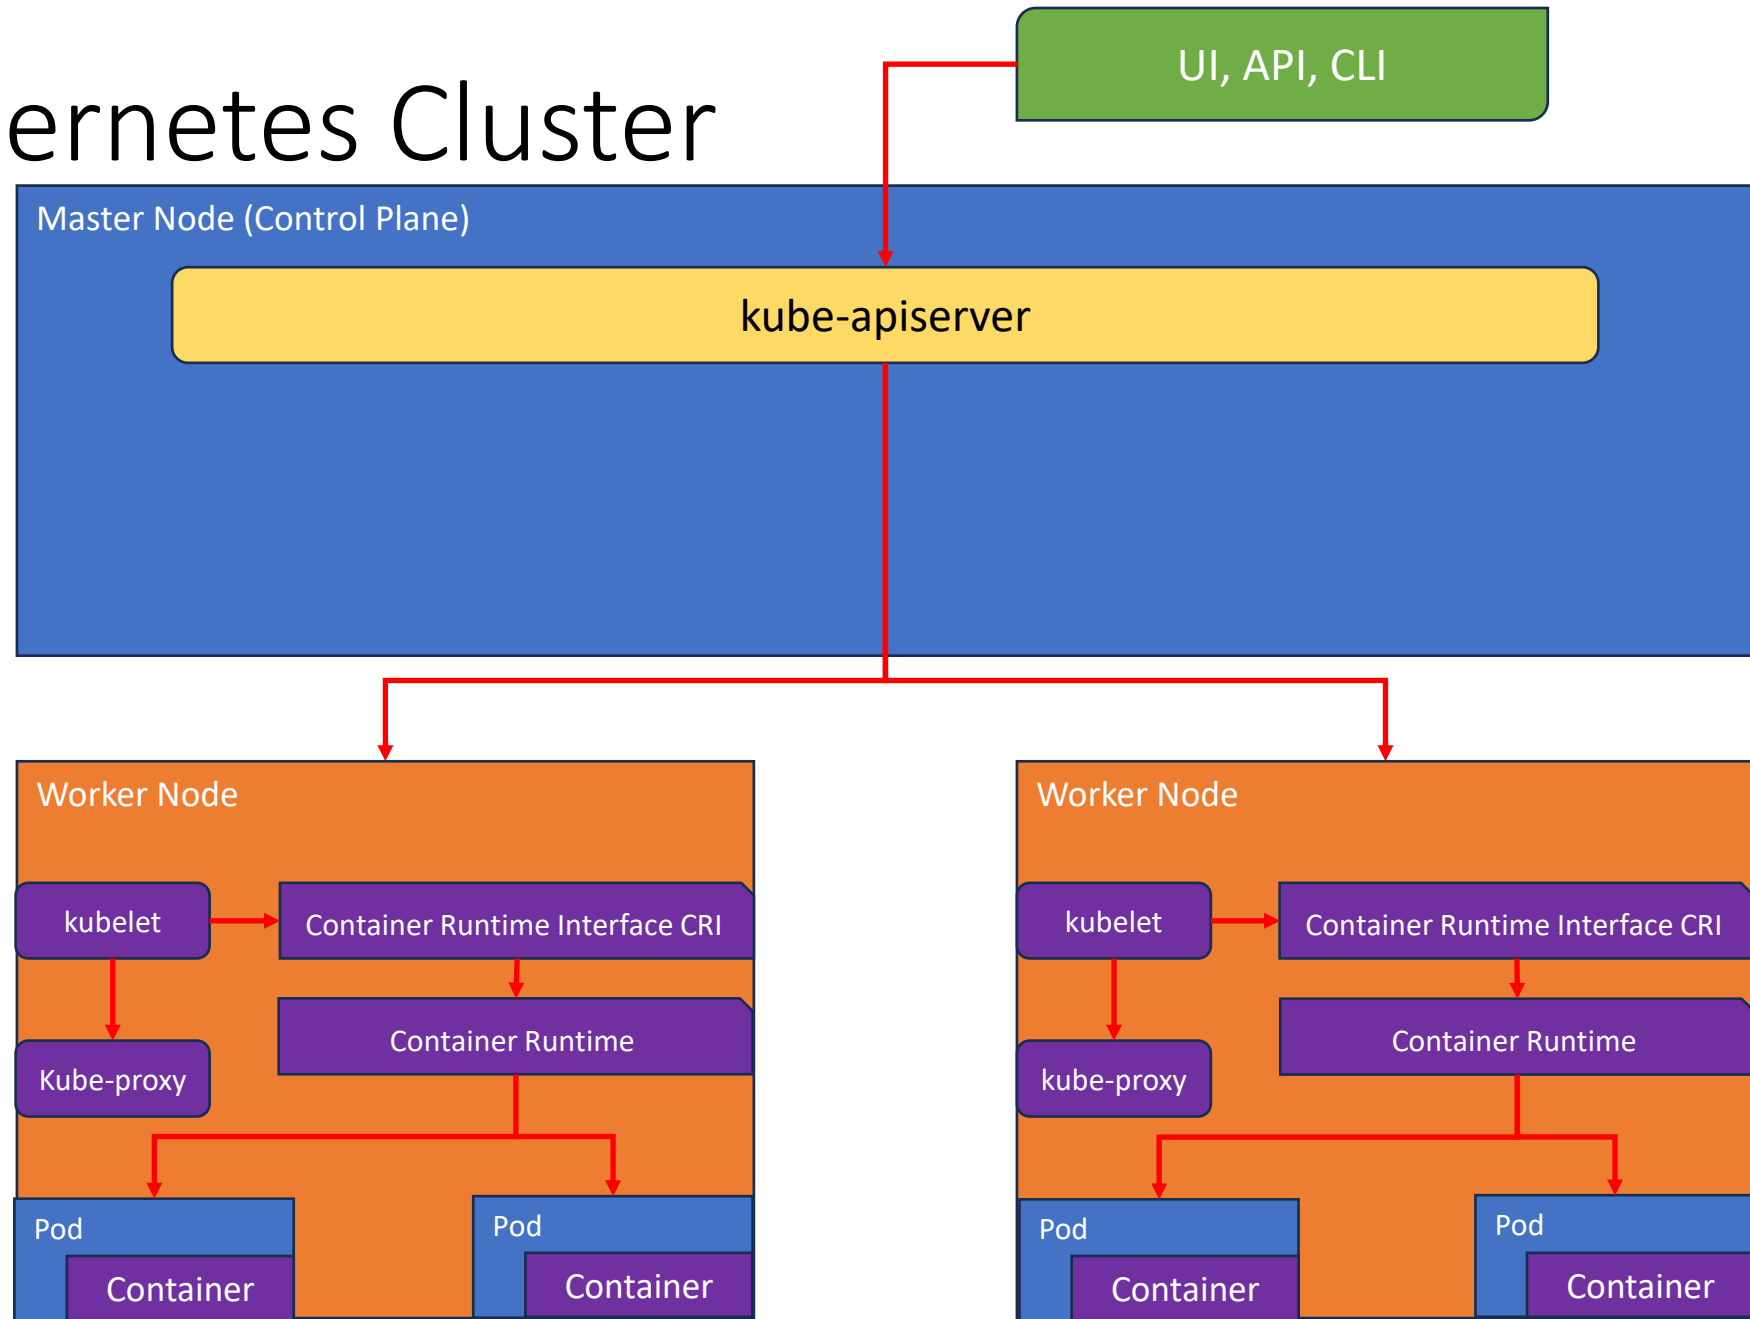
Kubernetes API via kubectl



Unit 1 – Kubernetes Cluster

API server – entry point to the k8s cluster

Exposes k8s API as RESTful API –
Can be accessed by UI, CLI (kubectl)
to control the cluster



Kubernetes API via kubectl



Kubectl syntax is as follows

kubectl [command] [TYPE] [Name] [flags]

kubectl get pods

kubectl get pods --namespace my-namespace

kubectl get services

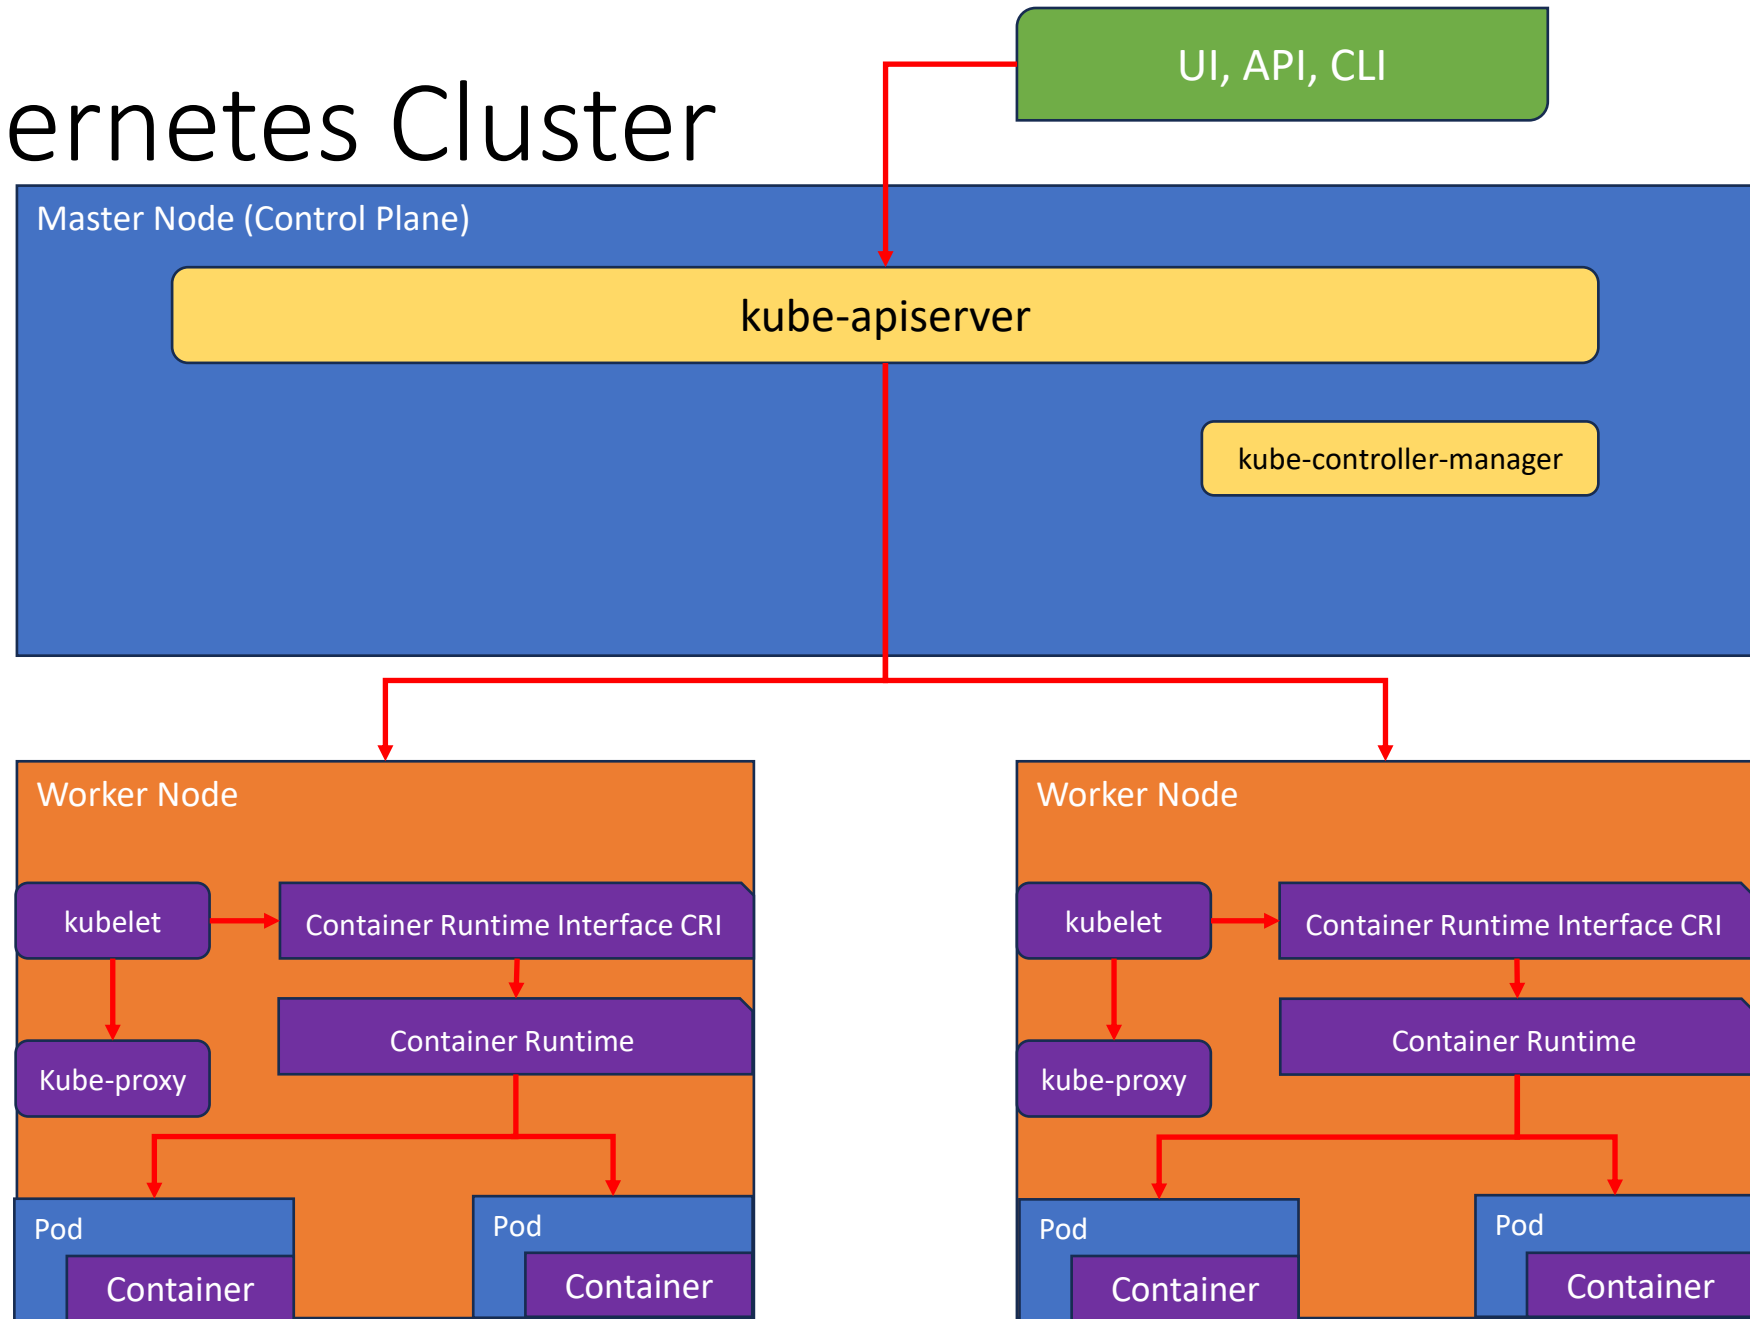
kubectl get replicaset

Unit 1 – Kubernetes Cluster

Controller-manager

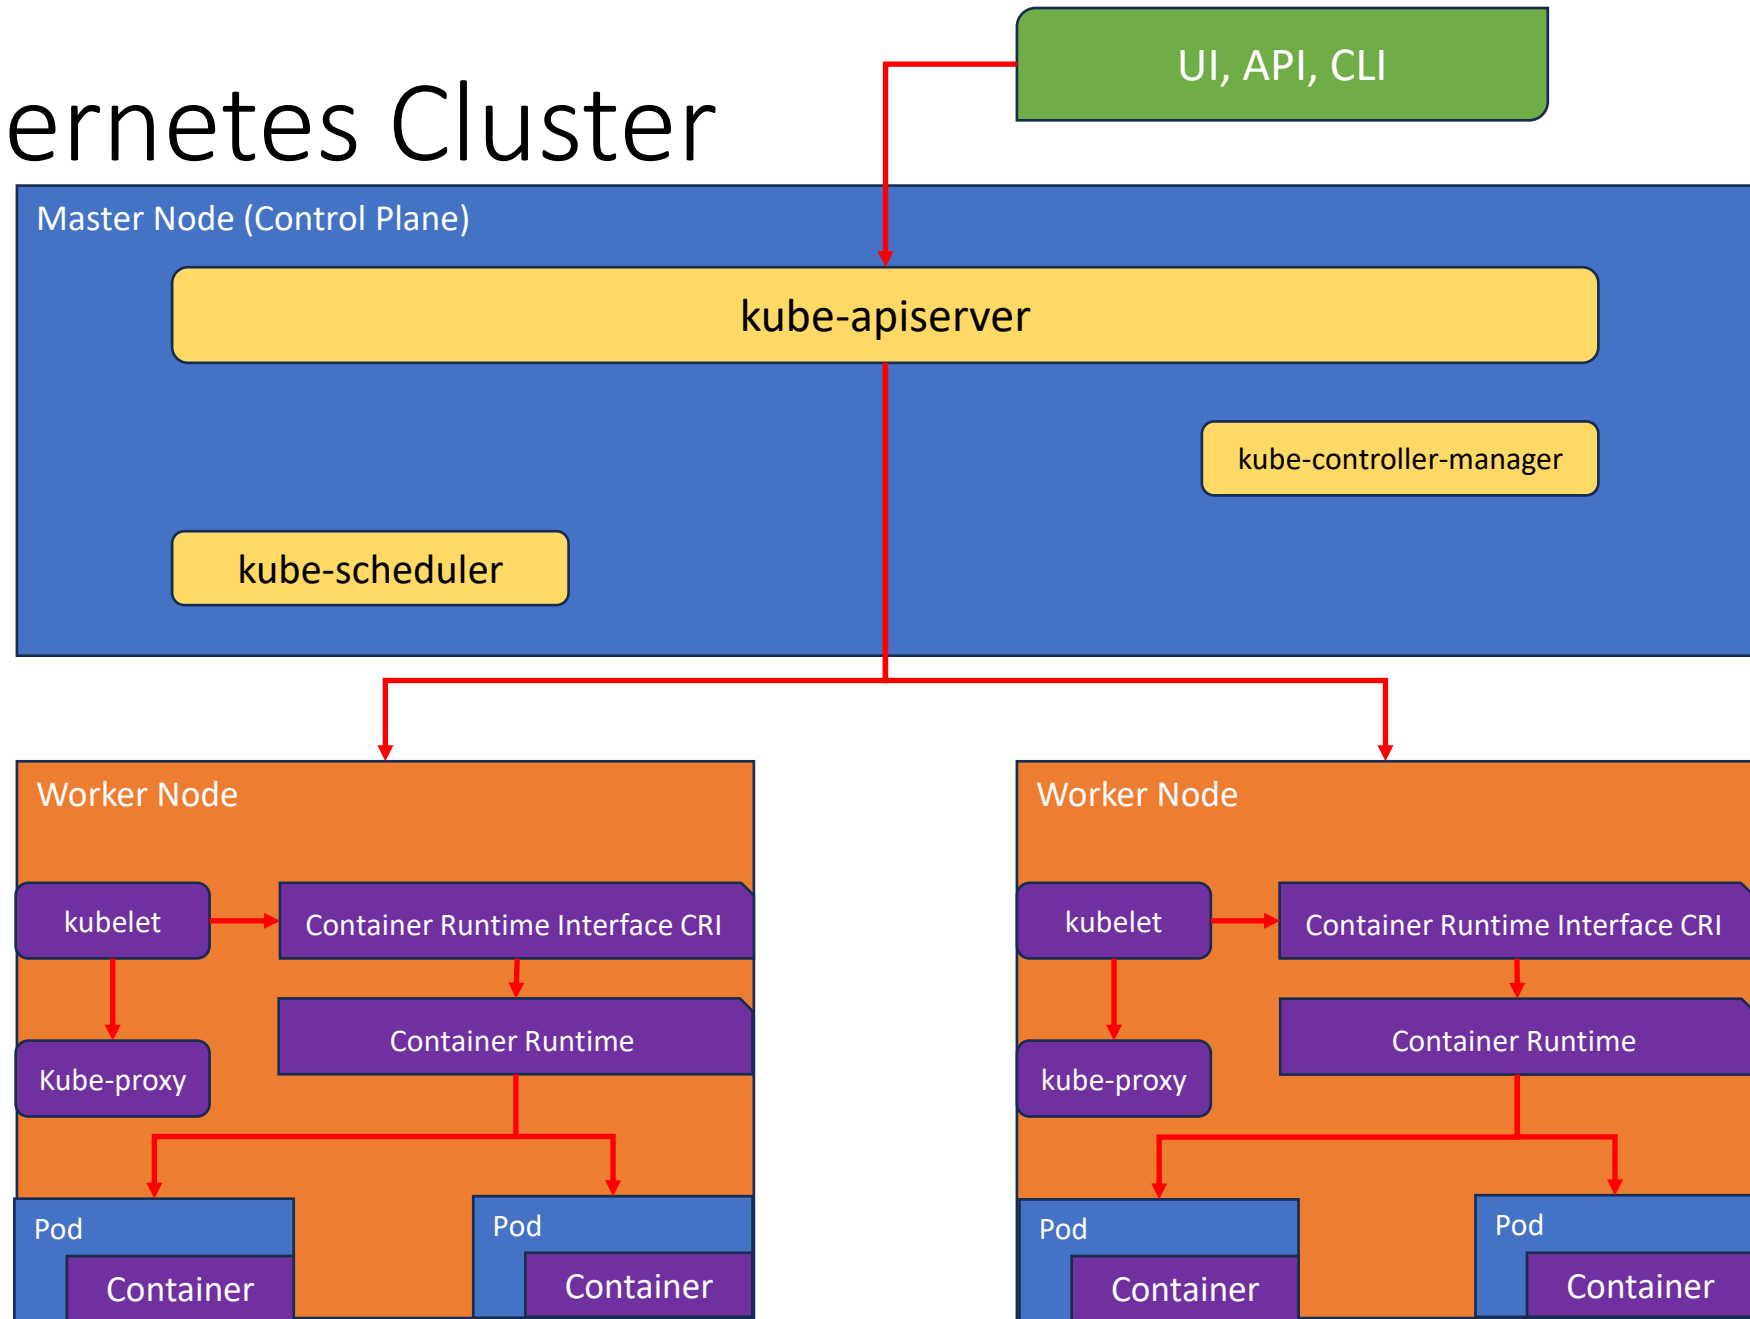
Responsible for noticing and responding when nodes go down

Responsible for maintaining the correct number of pods



Unit 1 – Kubernetes Cluster

Responsible for scheduling pods on worker nodes



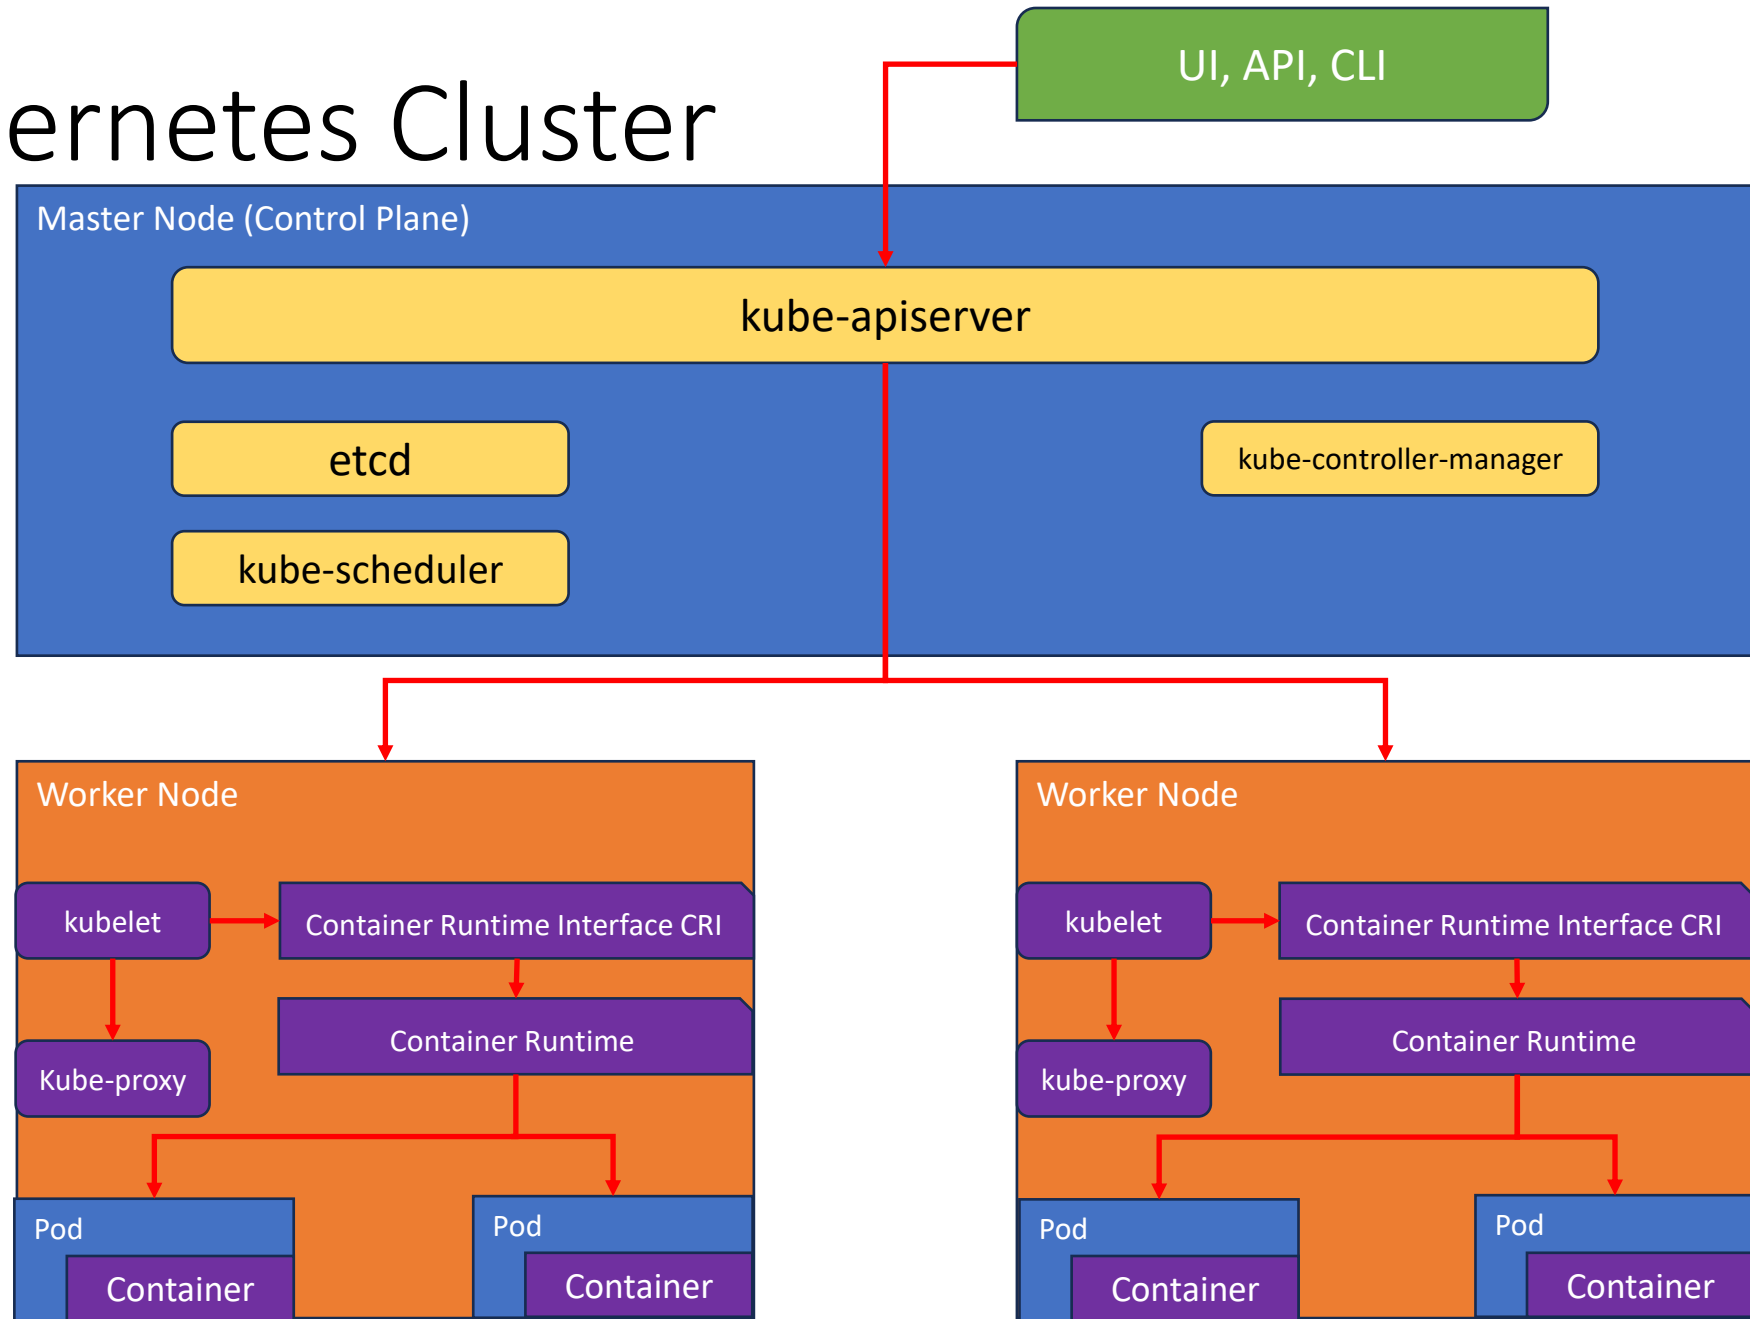
Unit 1 – Kubernetes Cluster

etcd is a key-value storage

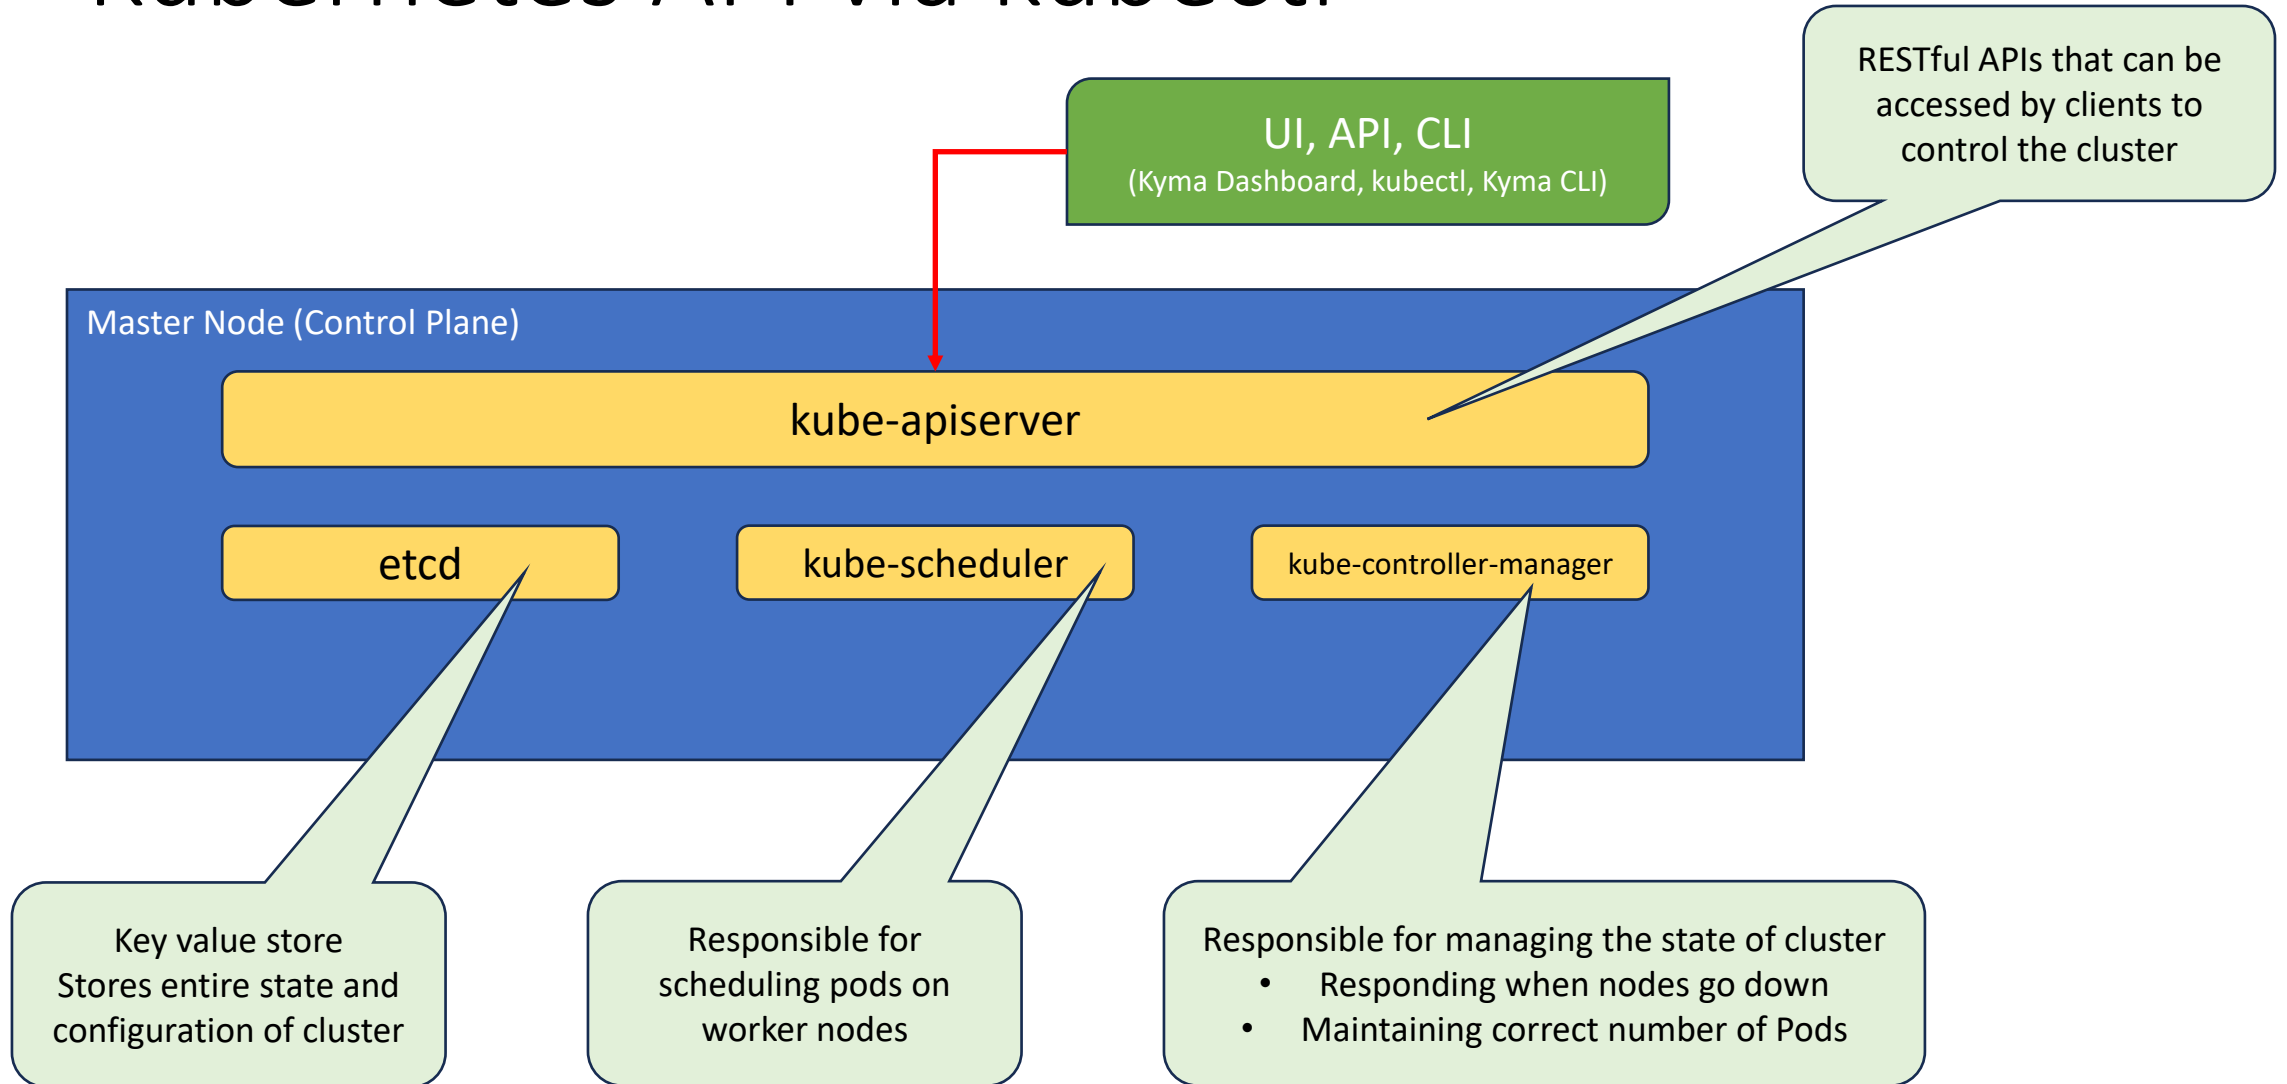
Contains configuration data. Holds the current info of k8s cluster

Master node is the heart of the k8s cluster

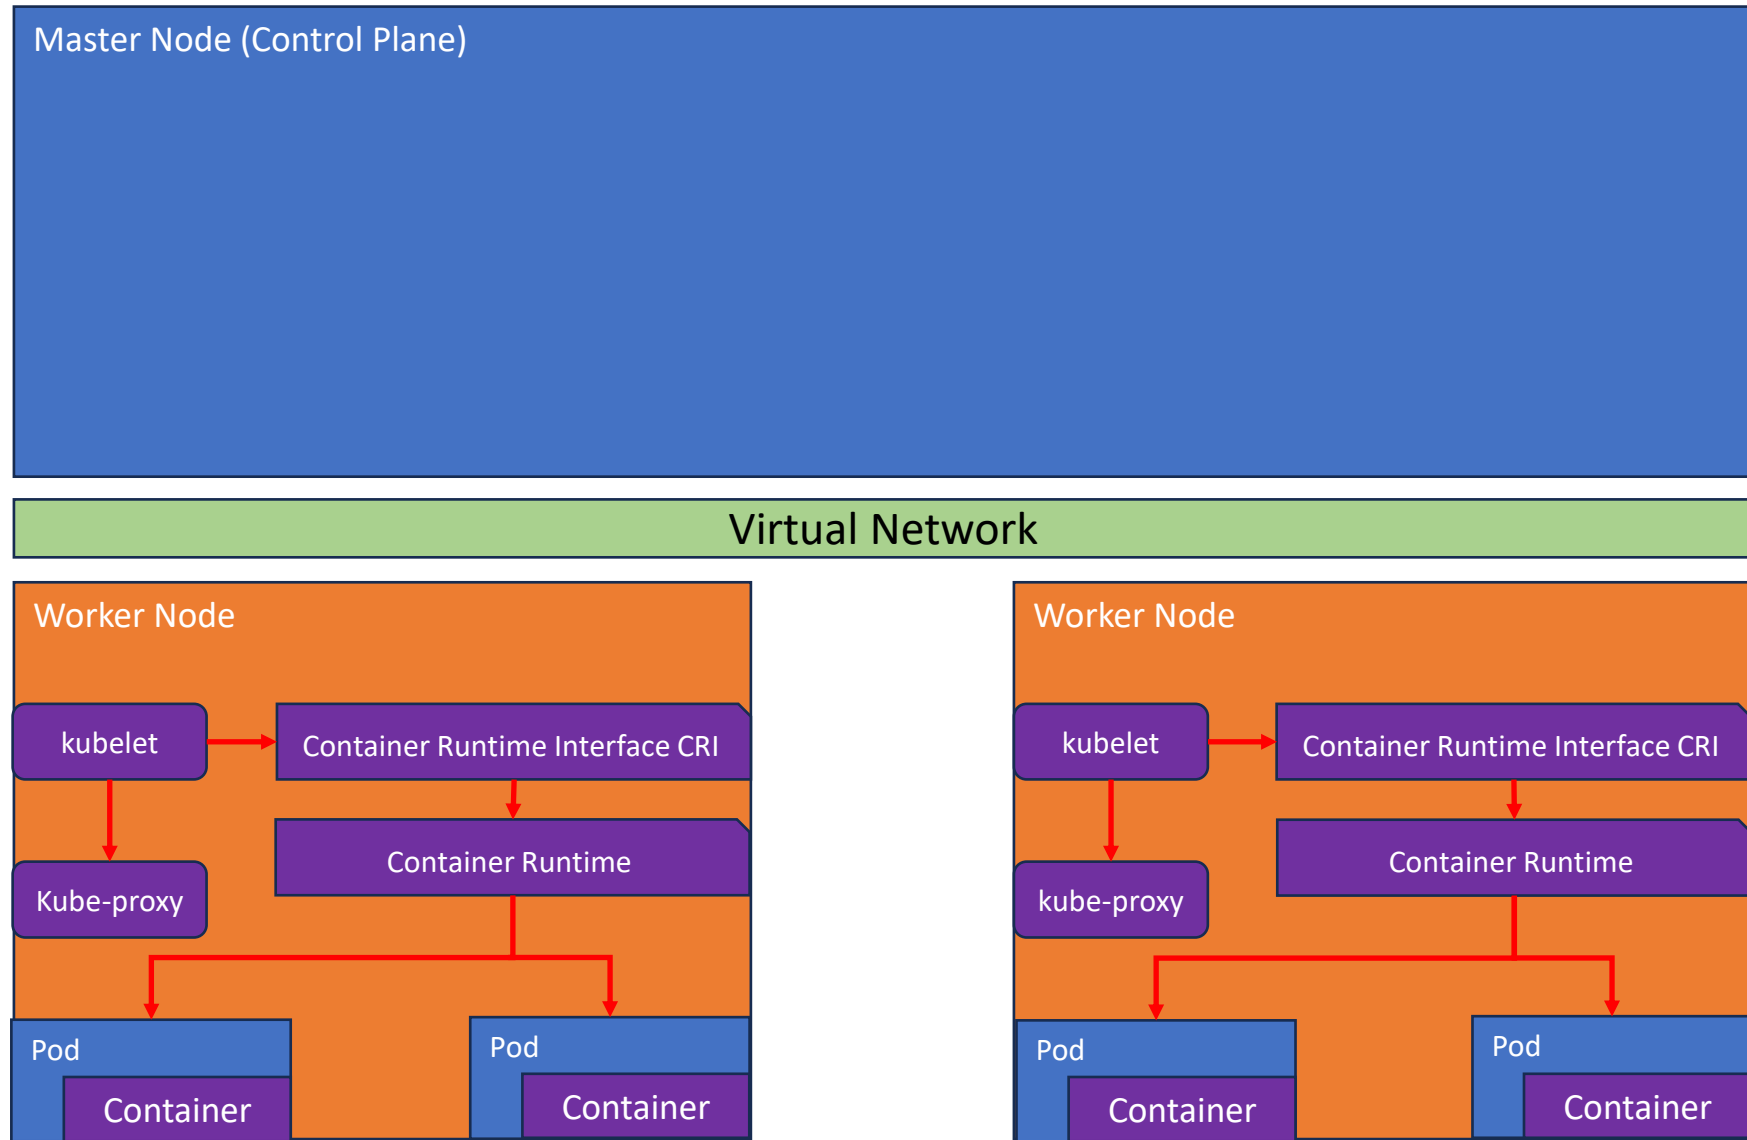
Can be used for backup and restore based on etcd snapshot



Kubernetes API via kubectl



Unit 1 – Kubernetes Cluster

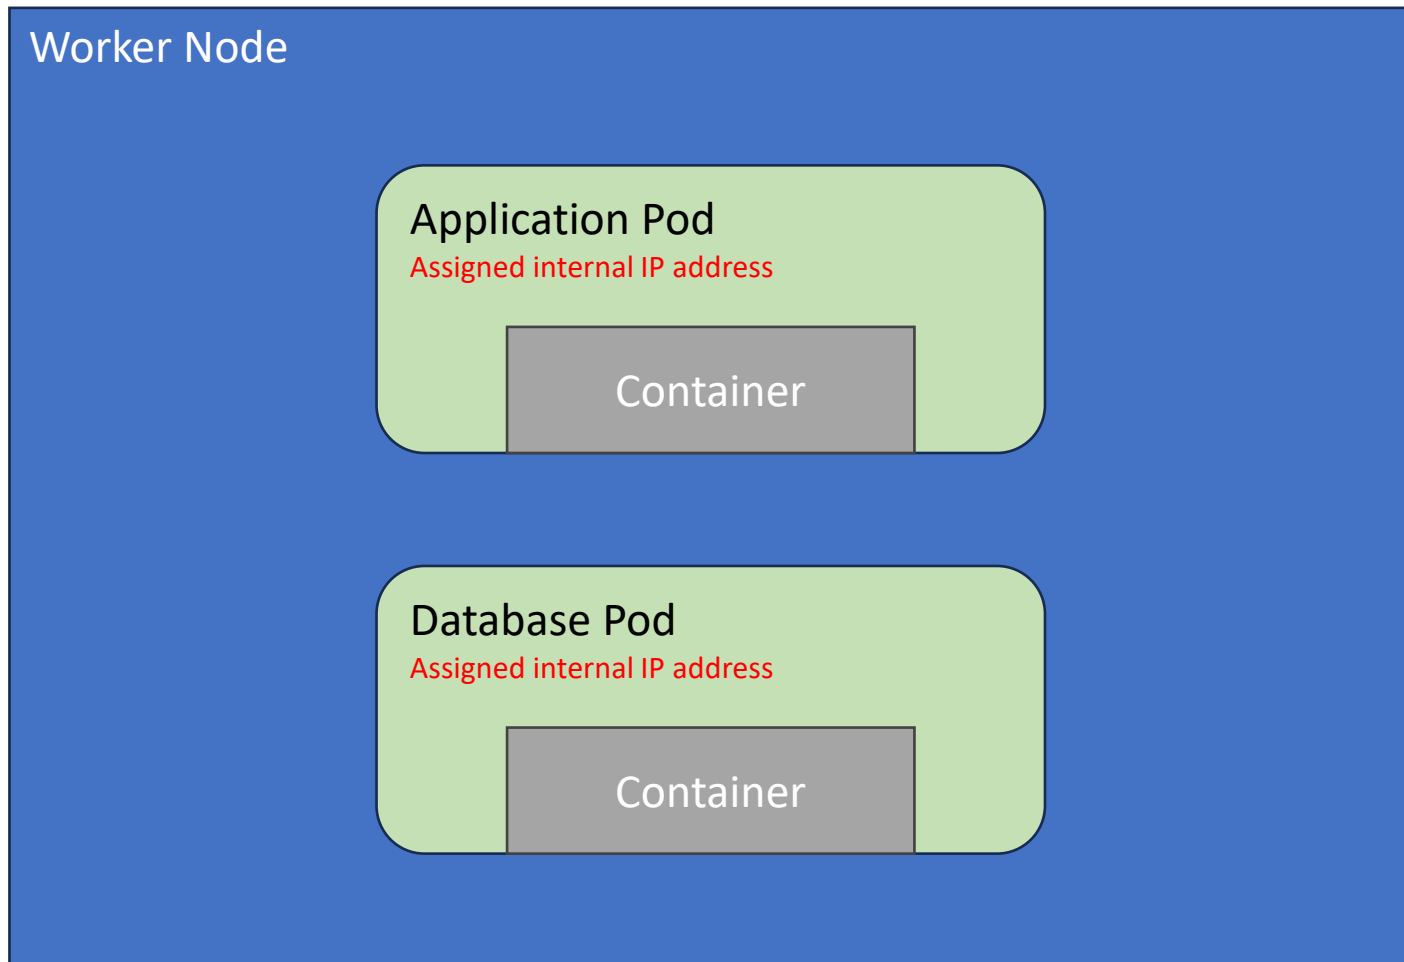


Unit 1 – Kubernetes Components

Each Pod is assigned an internal IP address

This allows **Application Pod** and **Database Pod** to talk to each other

Pods are ephemeral – can die easily



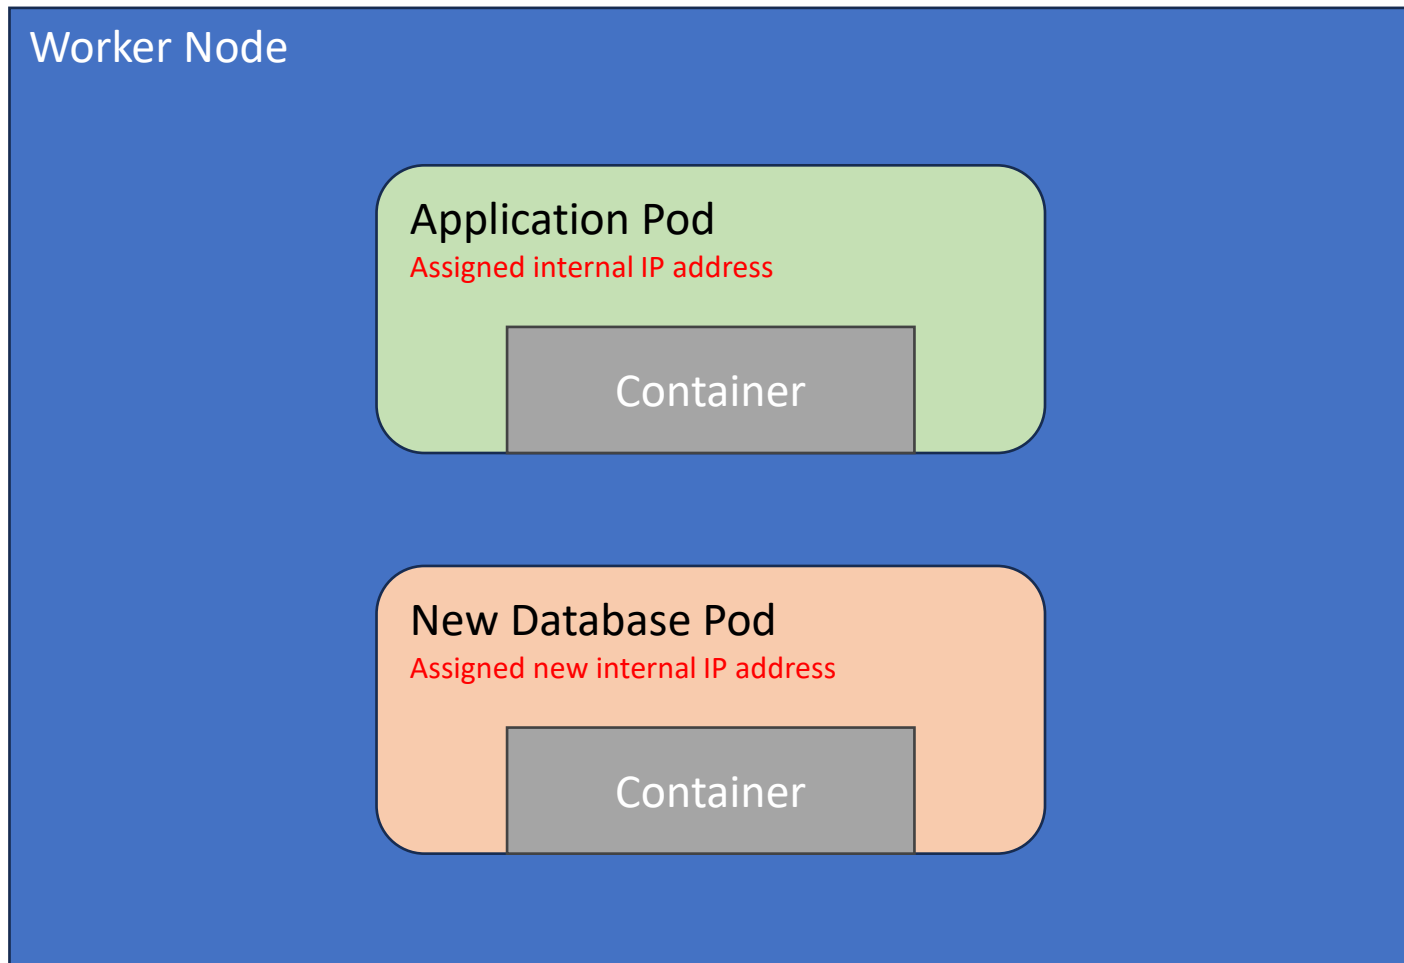
Unit 1 – Kubernetes Components

When a Pod dies, k8s will create a new Pod in the cluster

New Pod is **assigned a new internal IP address**

Oh no !! **Application Pod** cannot talk to the **New Database Pod**

Service component to the rescue

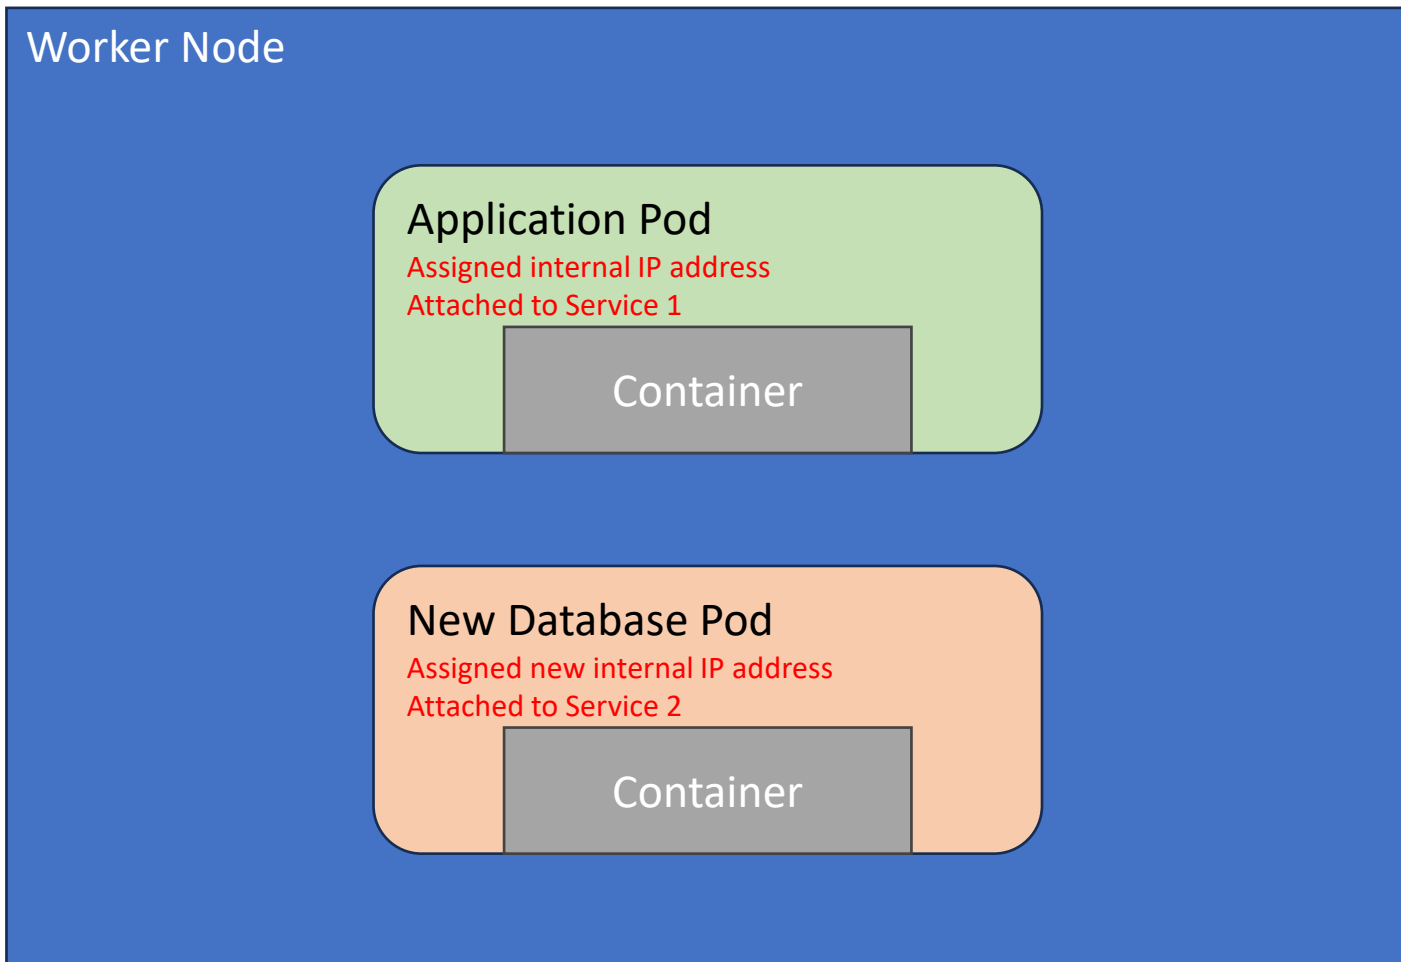


Unit 1 – Kubernetes Components

Service is a permanent IP address attached to each **Pod** – well, not too attached though

When the Pod dies, the **Service** does not die along with it. The lifecycle of the **Service** is independent of the lifecycle of the **Pod**

So now the **Application Pod** can communicate to the **New Database Pod** using Service (Permanent IP address)



Unit 1 – Kubernetes Components

Ok, this is all good. But how do I access this **Application Pod** from the outside world

You can create either an External Service or an Internal Service

Application Pod → External Service

New Database Pod → Internal Service

But I don't want to use an IP address to connect... I want a friendly name !!

Ingress component to the rescue

Worker Node

Application Pod

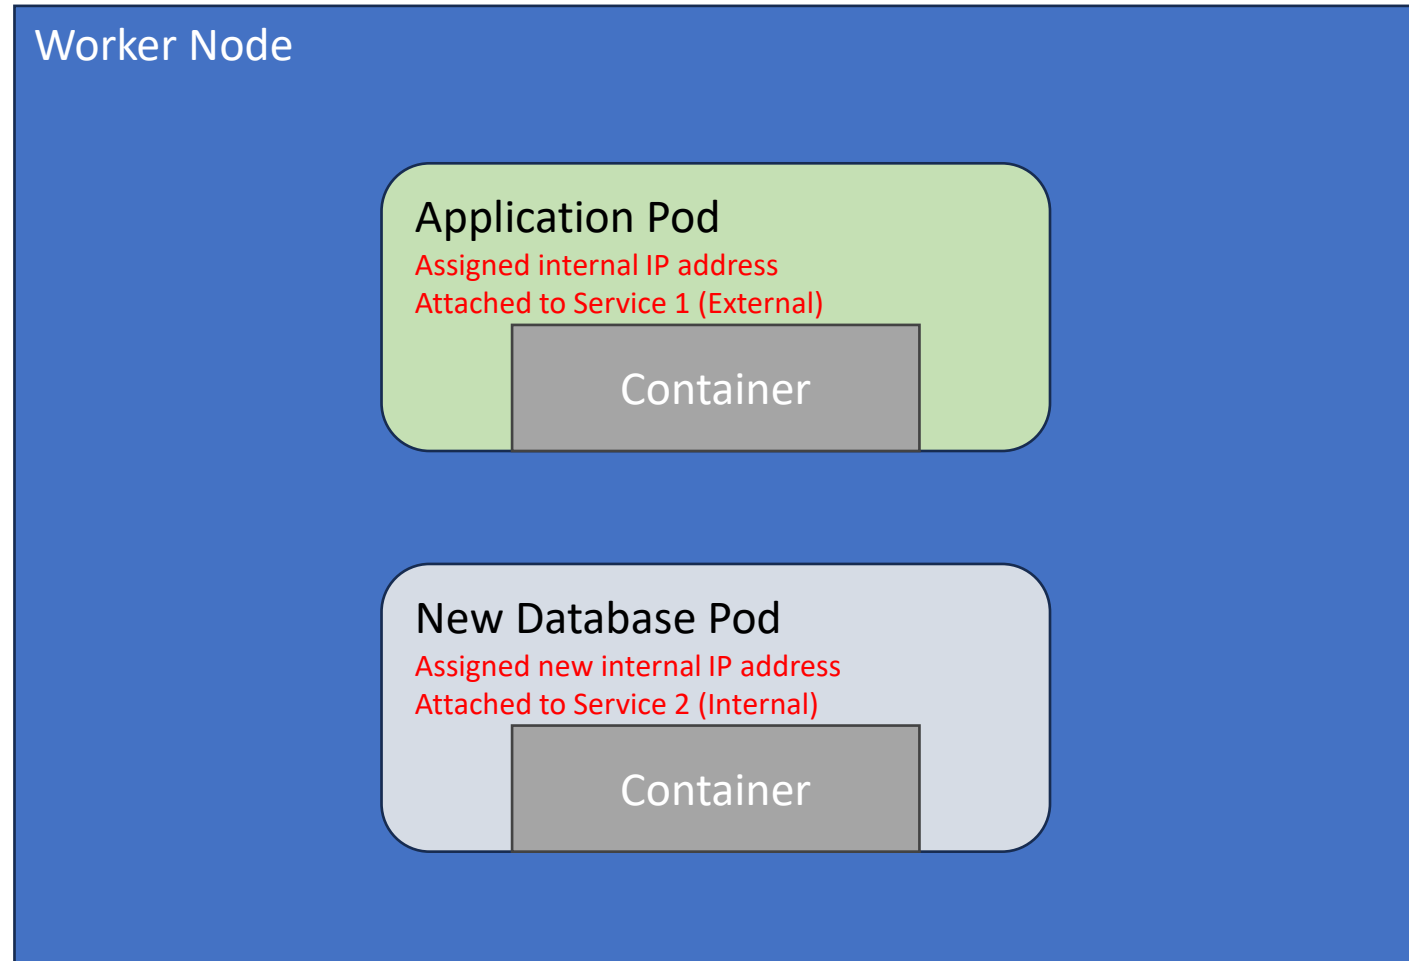
Assigned internal IP address
Attached to Service 1 (External)

Container

New Database Pod

Assigned new internal IP address
Attached to Service 2 (Internal)

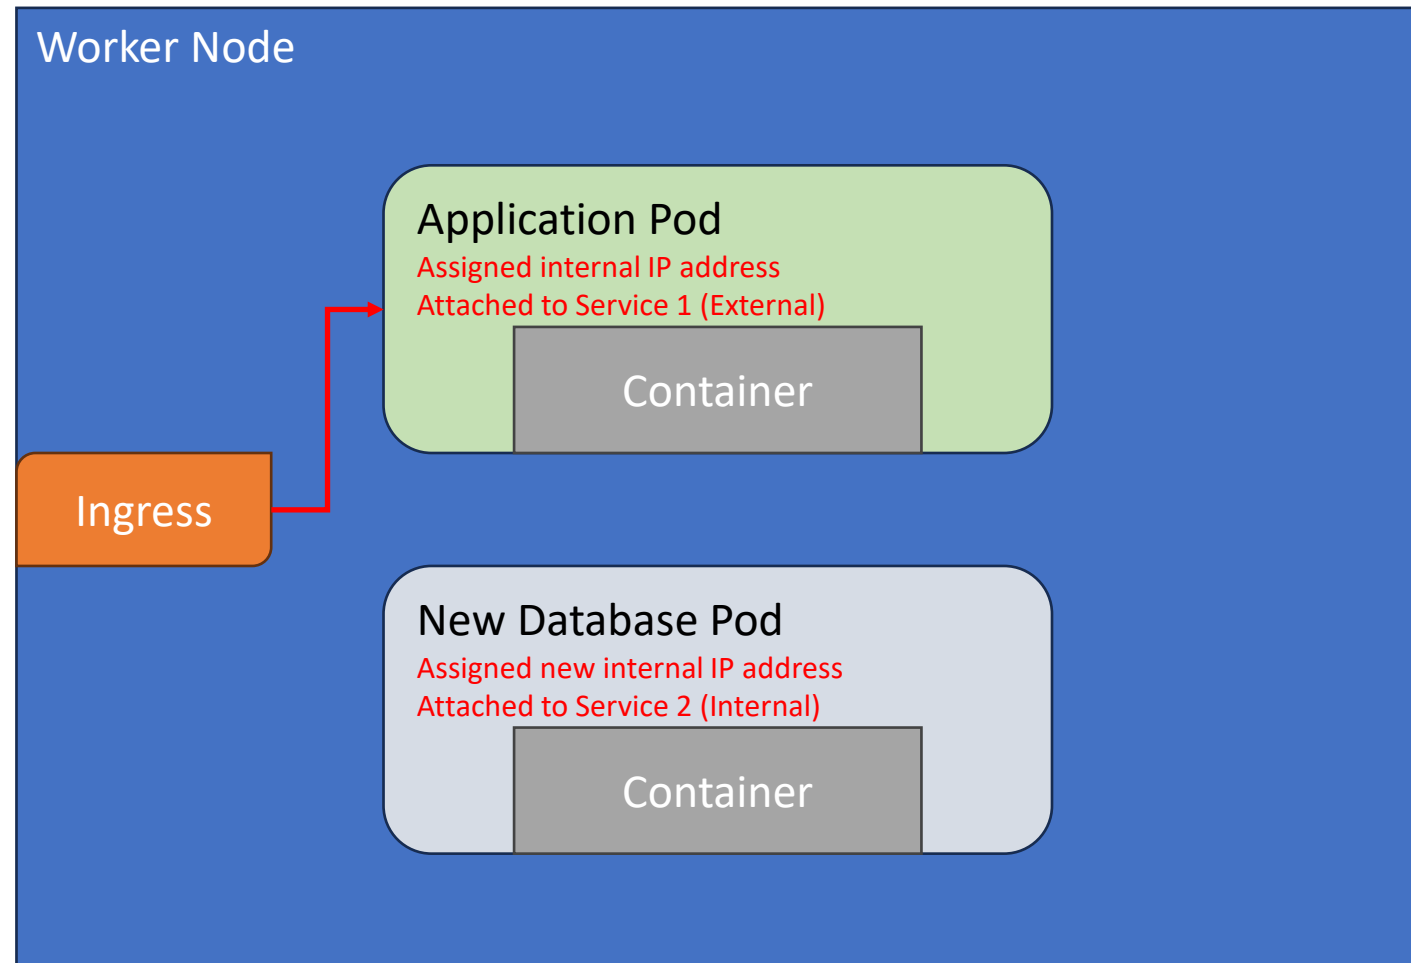
Container



Unit 1 – Kubernetes Components

So the request first goes to Ingress

Ingress routes the request to the right Service
(in this case, Service 1 of the Application Pod)



Unit 1 – Kubernetes Components

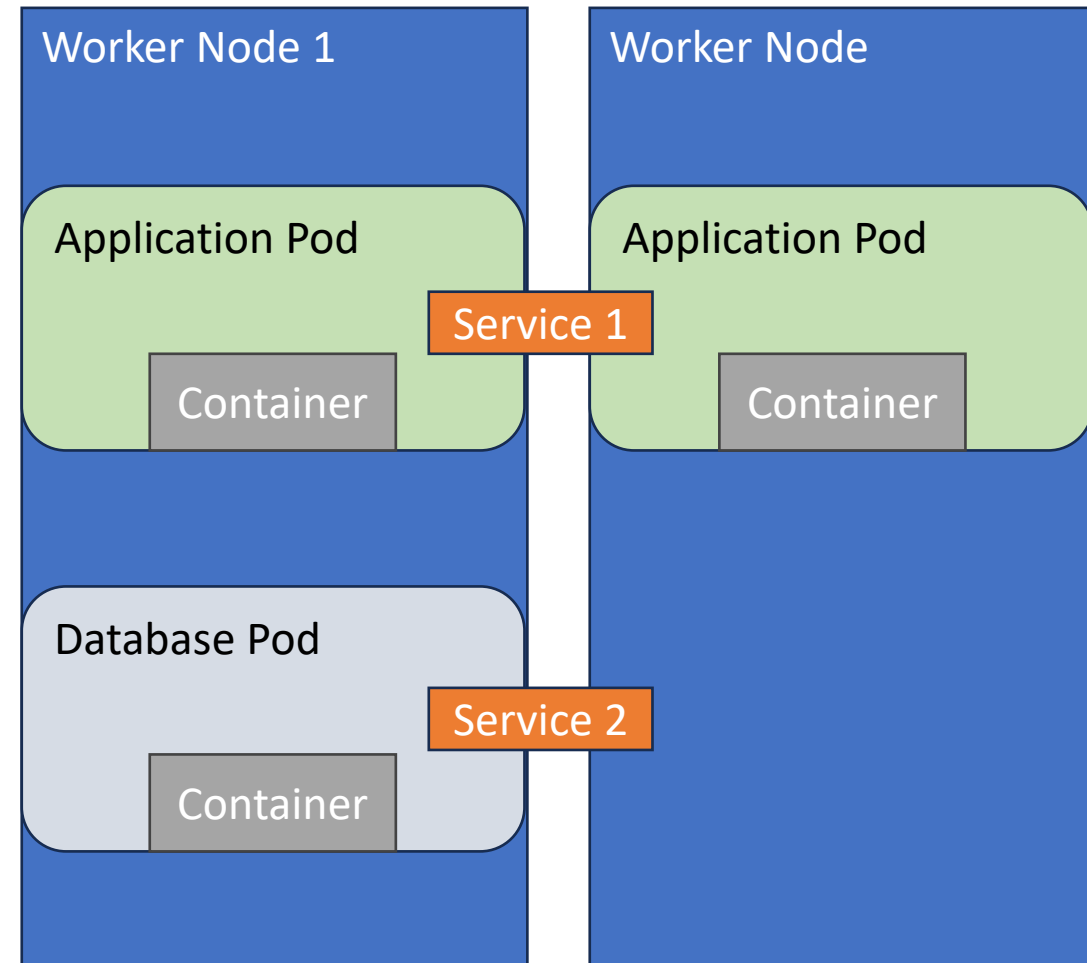
For high availability, we typically have more than 1 Application Pod

Do we need to define multiple Application Pods in the code ? **No**

You only have a single definition of Application Pod

But you mention that you want x number of copies of that

[ReplicaSets](#) define how many Pods we need in the cluster



Unit 1 – Kubernetes Components

For high availability, we typically have more than 1 Application Pod

Do we need to define multiple Application Pods in the code ? **No**

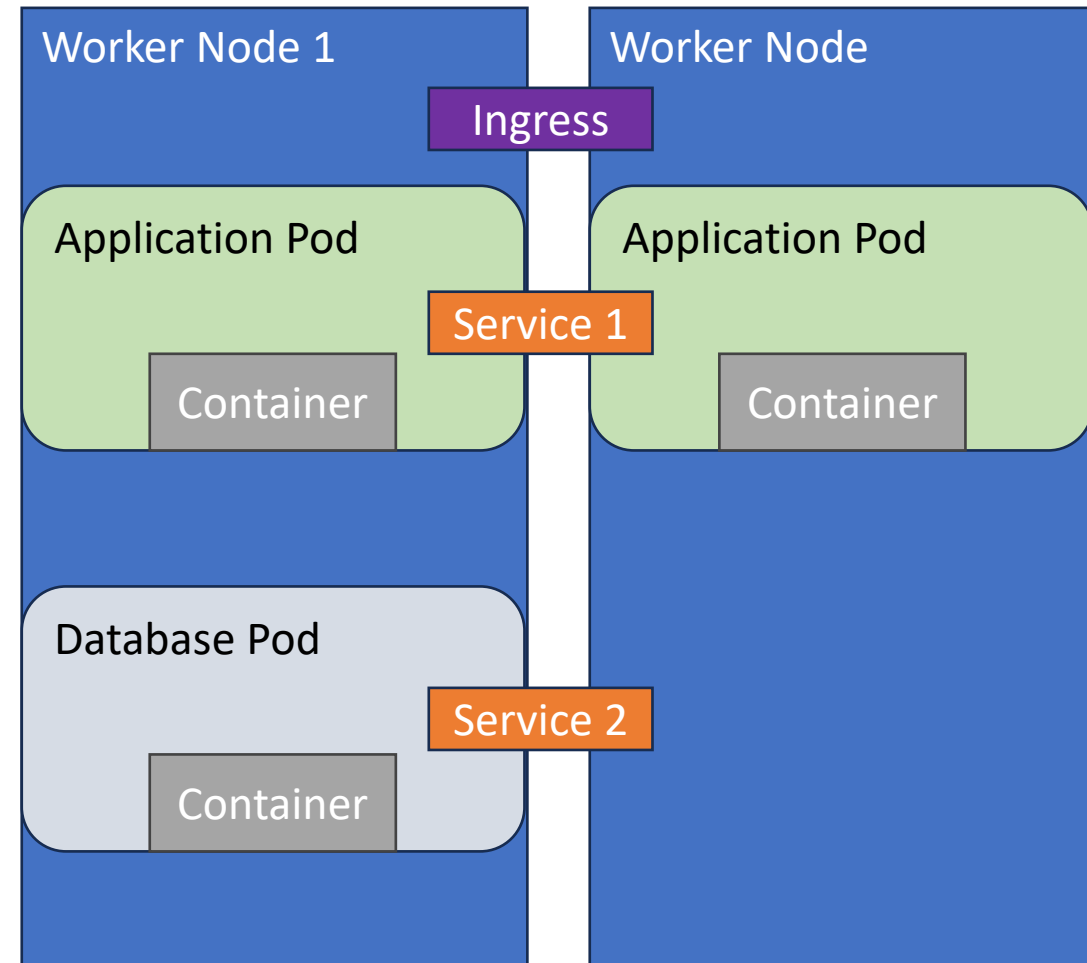
ReplicaSets define how many Pods we need in the cluster

Service can also act as a load balancer

ReplicaSets are often indirectly used through Deployments which provides additional features like rolling updates and rollback

Layers of abstraction

- Deployment manages a ReplicaSet
- ReplicaSet manages a Pod
- Pod is an abstraction of a Container



Unit 2 – Discovering Kyma

- Kyma is an open-source project
- Makes it easier to develop and run applications on Kubernetes
- Features provided by Kyma (extending functionality of k8s)
 - Serverless
 - Eventing
 - Observability
 - Service Integrations
 - Service Mesh
 - API Gateway

Unit 2 – Benefits of SAP BTP Kyma Runtime

- Fully managed Kubernetes cluster
- Latest features of Kubernetes and its ecosystem
- Additional features through Kyma (Serverless, Eventing etc.)
- Integration with other SAP services and products
- Elimination of operations' overhead

Enabling SAP BTP, Kyma Runtime



Kyma Dashboard



Kubectl for Kyma




Unit 3 – Discovering Kubernetes Workloads

Deployment

- Declaratively define the blueprint for Application Pod
- How many copies
- You would mostly work with deployments

StatefulSet

- Database cannot be replicated via deployment
- StatefulSet used to manage stateful applications (Database)
- To prevent any inconsistencies
- More complex than working with deployments



Serverless Functions

The screenshot displays the Kyma dashboard interface. The left sidebar contains a navigation menu with the following items: Overview, Events, Workloads (expanded), Functions (highlighted with a red circle and the number 1), Deployments, Stateful Sets, Daemon Sets, Cron Jobs, Jobs, Replica Sets, Pods, Discovery and Network (expanded), API Rules (highlighted with a red circle and the number 2), Ingresses, Services, Horizontal Pod Autoscalers, Network Policies, Istio, Service Management, Storage, Apps, and Configuration. The main content area shows the 'default' namespace details, including labels, creation time, and status (ACTIVE). Below this, there are sections for 'Healthy Resources' (Pods and Deployments, both showing 0/0), 'Resource Consumption' (Memory Requests and Memory Limits, both showing 0/0), a CPU usage graph (showing a single data point at 18:37), a 'Limit Range' table (with no entries found), and a 'Resource Quota' section.

Functions 1

API Rules 2 Virtual service is created by default

default

Labels: istio-injection=enabled, kubernetes.io/metadata.name=default

Created: 17 days ago

Status: ACTIVE

Buttons: Edit, Deploy new workload +, Upload YAML +, Delete

Healthy Resources

Pods: 0/0, Deployments: 0/0

Resource Consumption


Memory Requests: 0/0, Memory Limits: 0/0

CPU usage graph: 1h, 3h, 6h

Limit Range table: No entries found

Resource Quota section

Serverless Functions



Deployments

Kyma default ▾

Back to Cluster Details

Overview

Events

Workloads ▾

- Functions
- Deployments **1**
- Stateful Sets
- Daemon Sets
- Cron Jobs
- Jobs
- Replica Sets
- Pods

Discovery and Network ▾

- API Rules **3**
- Ingresses
- Services **2**
- Horizontal Pod Autoscalers
- Network Policies

Istio >

Service Management >

Storage >

Apps >

Configuration >

Labels:

Created: 17 days ago Status: **ACTIVE**

Edit **Deploy new workload +** **Upload YAML +** **Delete**

Healthy Resources

Pods	Deployments
0/0	0/0

Resource Consumption

Memory Requests	Memory Limits
0/0	0/0

CPU ▾ 1h 3h 6h

1

1

0 18:07 18:22 18:37 18:52 19:06 CPU

Limit Range

Create Limit Range +

Name	Created	Labels	Max	Min	Default	Default Request
No entries found						

Resource Quota

Create Resource Quota +

Creating and Managing a Deployment



Unit 3 – What is Helm ?

- Package manager for k8s
- Let's say you want to deploy gitlab
 - Might need the following...
 - Service to expose the endpoint
 - Persistent volume
 - Secret, ConfigMaps etc.
- Helm makes it super easy
 - *helm install gitlab*
 - *helm upgrade gitlab*
 - *helm uninstall gitlab*



gitlab

GitLab GitLab

GitLab is the most comprehensive AI-powered DevSecOps Platform.

Integration and delivery

★ 196 Helm chart

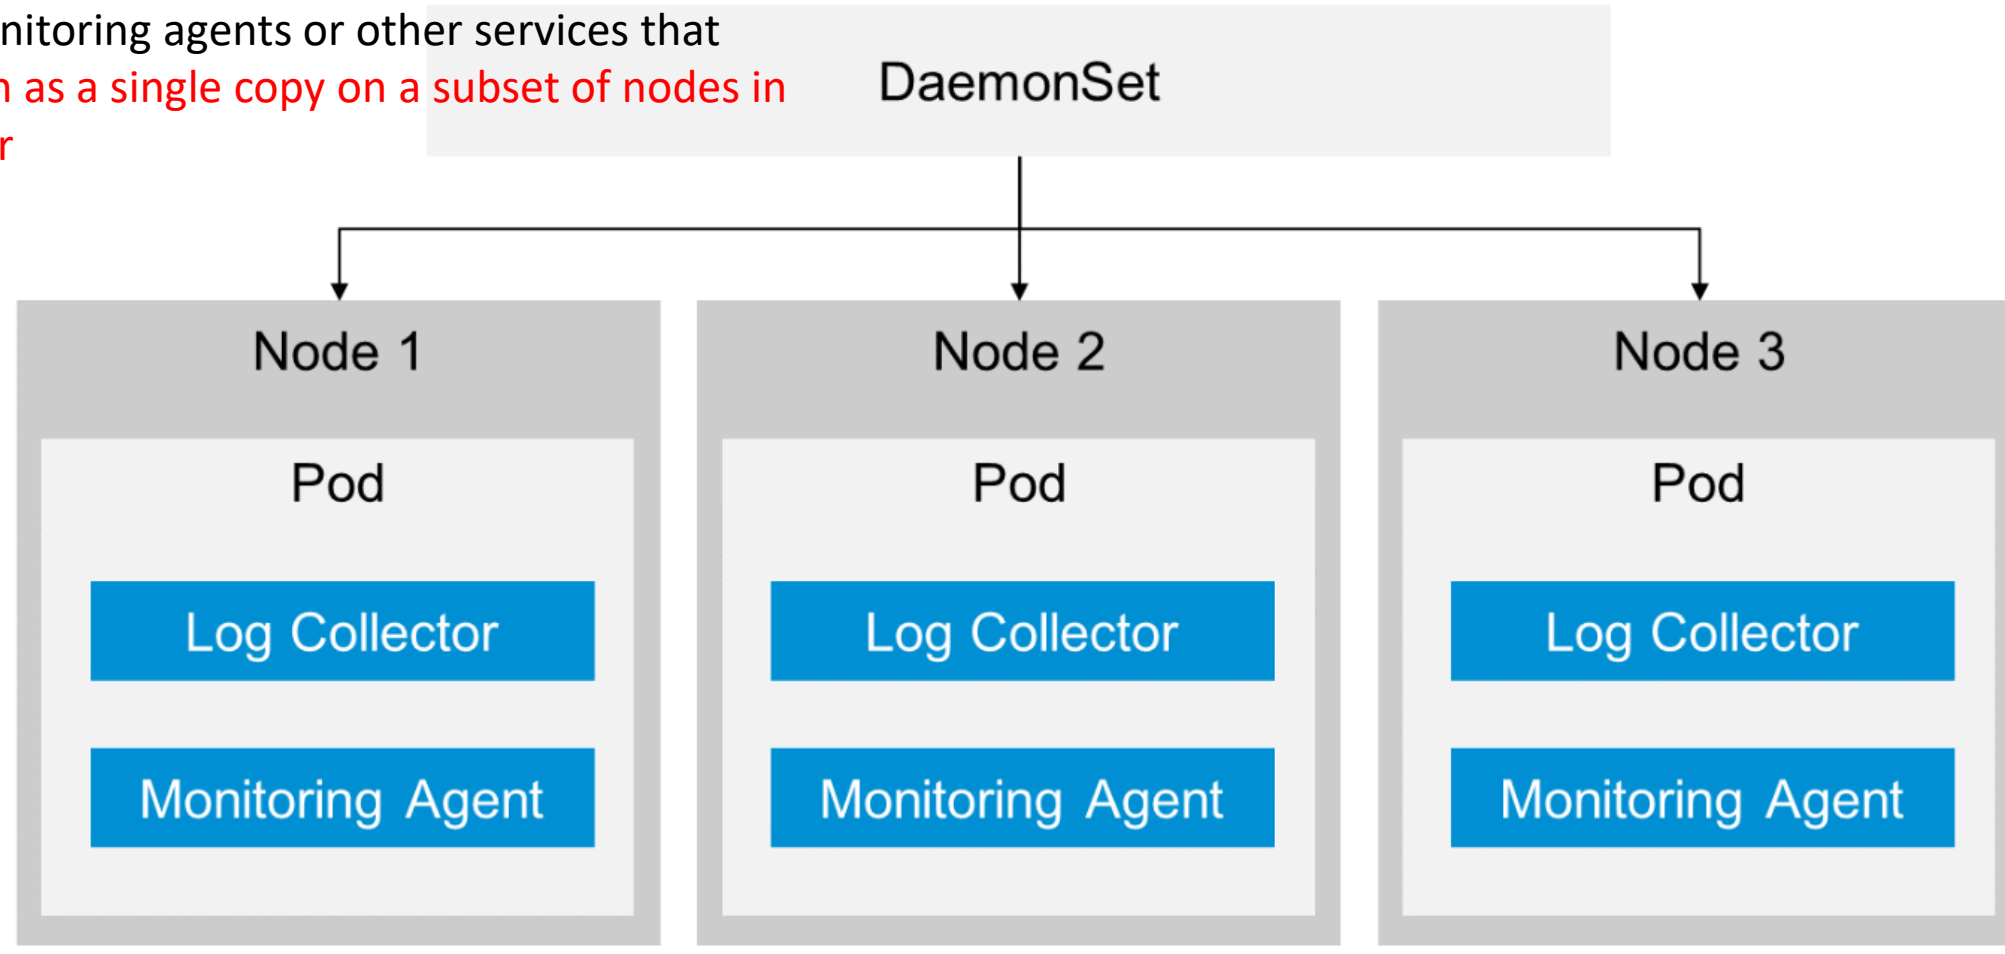
Updated 4 days ago

Version 7.9.2



Unit 3 – DaemonSets in Kubernetes

DaemonSets are helpful for services such as logging agents, monitoring agents or other services that need to **run as a single copy on a subset of nodes in your cluster**



Unit 3 – Jobs in Kubernetes



One-off / Non-parallel Job

Description

A job that runs once and is completed after successful termination

When to use

For example, a database migration



Parallel Job

Description

A job that runs multiple pods in parallel and is completed after all pods have terminated

When to use

For example, a batch job that needs to process a large amount of data

Jobs are generally used to run a workload only once

For example,

- Database migration job
- Some batch job

Supposed to run to completion and then terminate

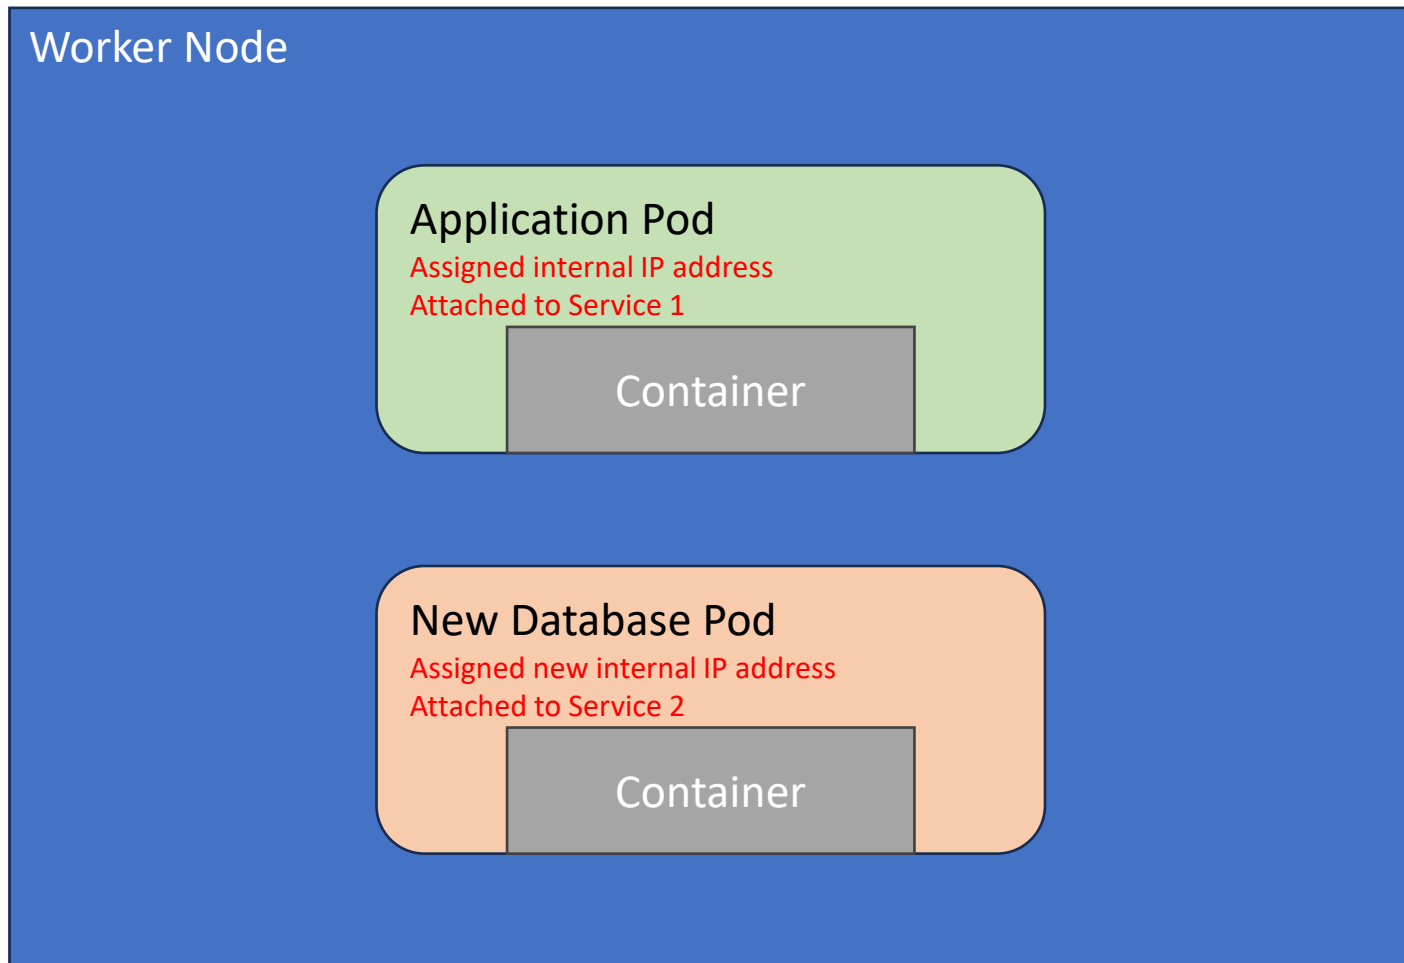
If a job fails, it will be restarted until it succeeds

Unit 4 – Discovering Services in Kubernetes

Service is a permanent IP address attached to each **Pod** – well, not too attached though

When the Pod dies, the **Service** does not die along with it. The lifecycle of the **Service** is independent of the lifecycle of the **Pod**

So now the **Application Pod** can communicate to the **New Database Pod** using Service (Permanent IP address)



Unit 4 – Discovering Services in Kubernetes

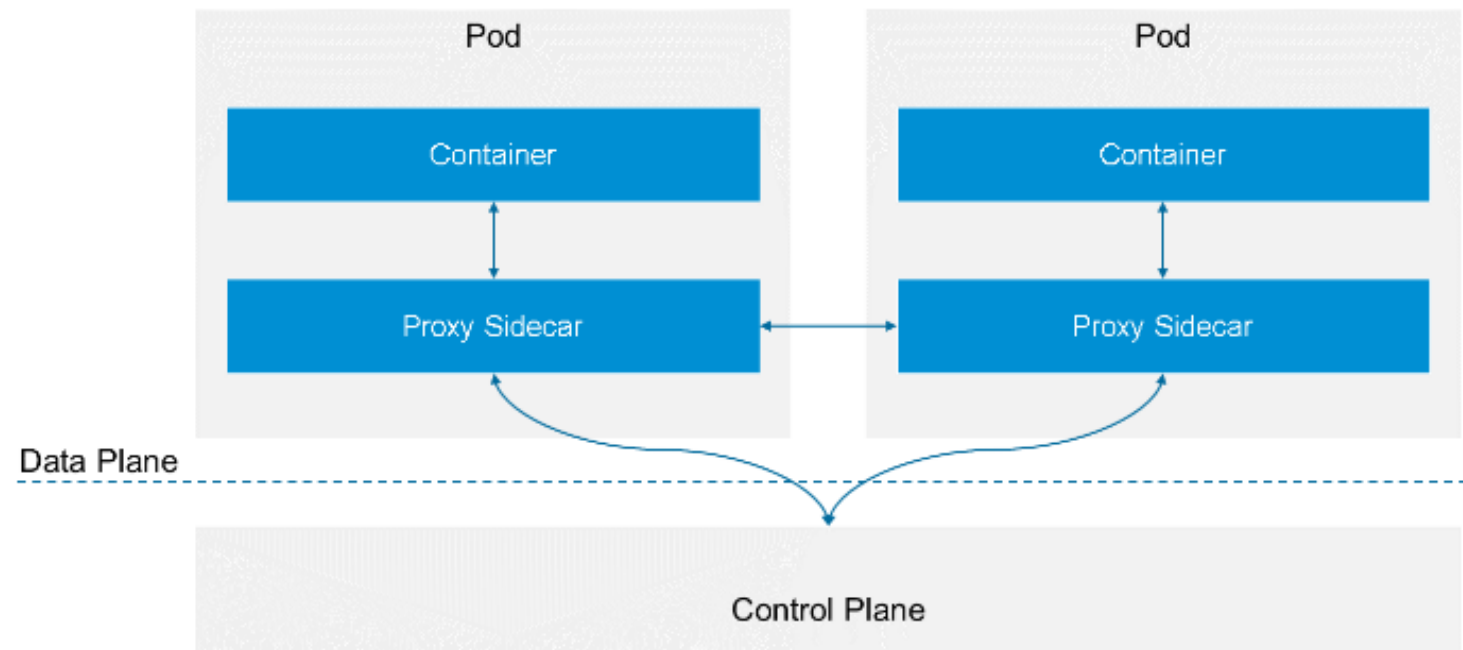
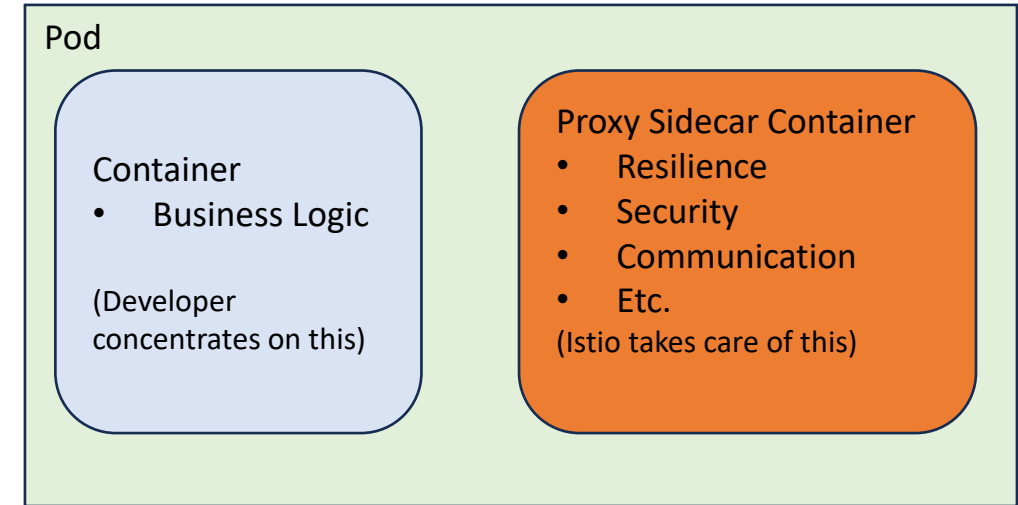


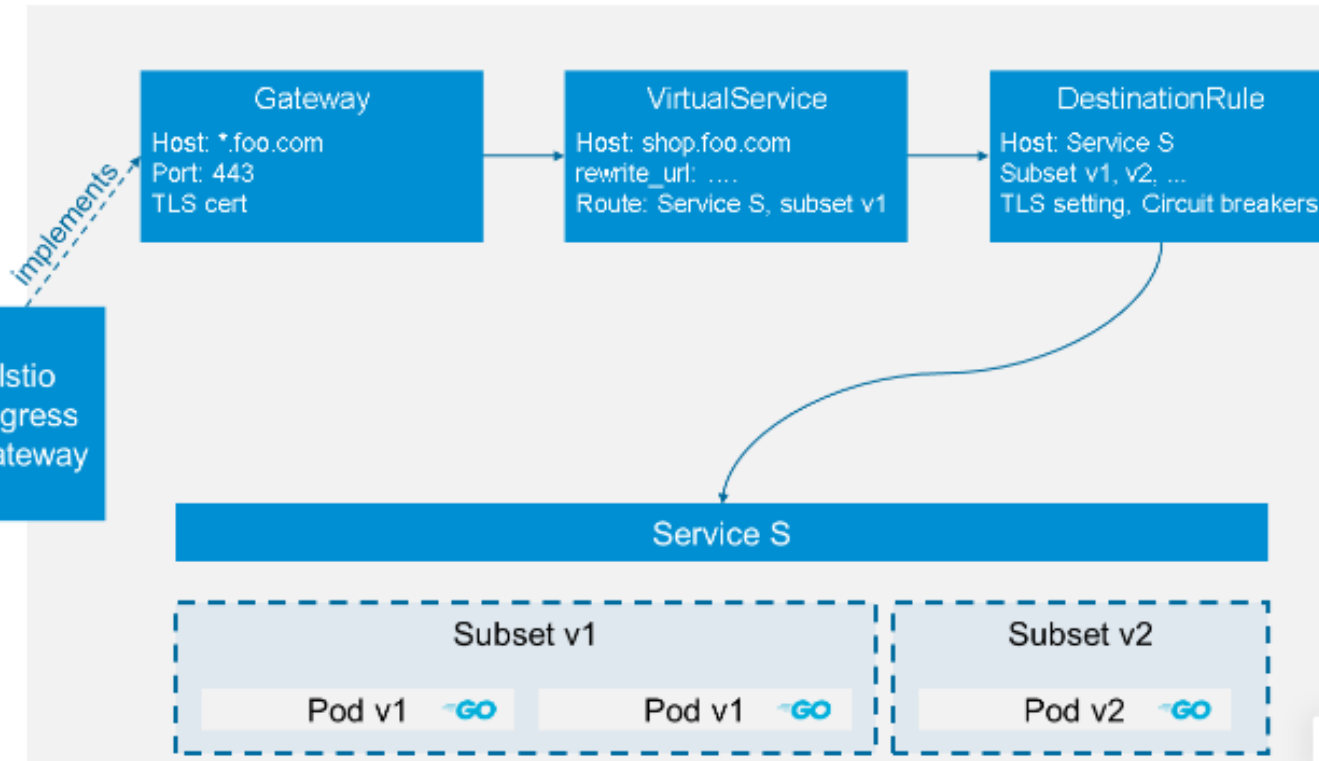
Unit 4 – Using API Gateway to expose Services

- API Gateway is used on top of k8s services to make it publicly available
- API Gateway is a custom configured Istio Ingress Gateway
- By creating an API Rule for a service
 - Creates an Istio Virtualservice
 - Create corresponding Oathkeeper Access Rules

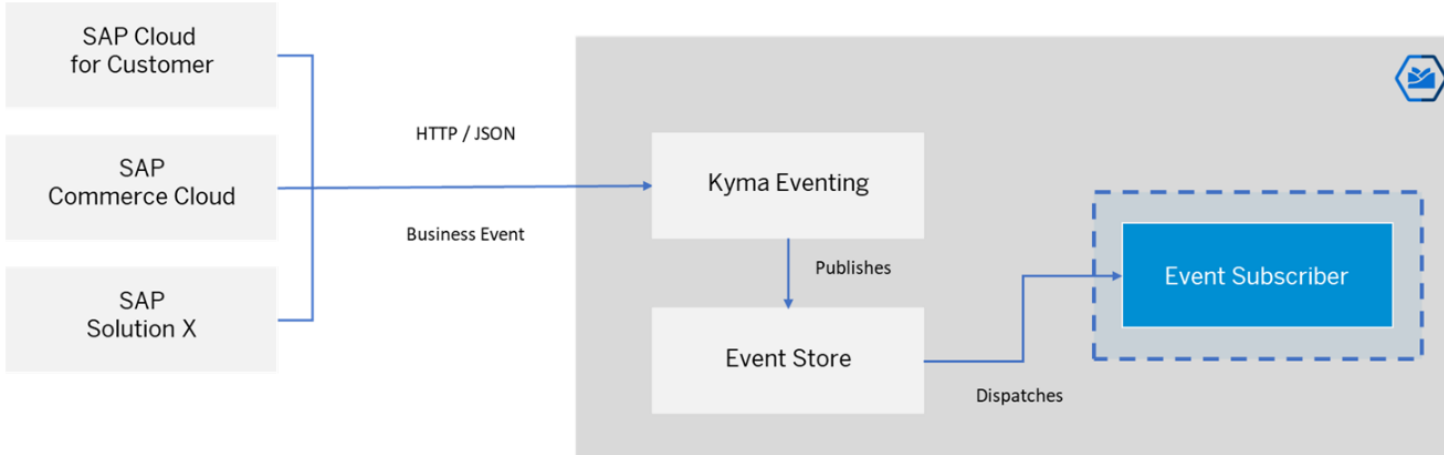
Unit 5 – Service Mesh

- Benefits of Service Mesh
 - Security
 - Observability
 - Traffic Management
 - Resilience
- Service Mesh Architecture



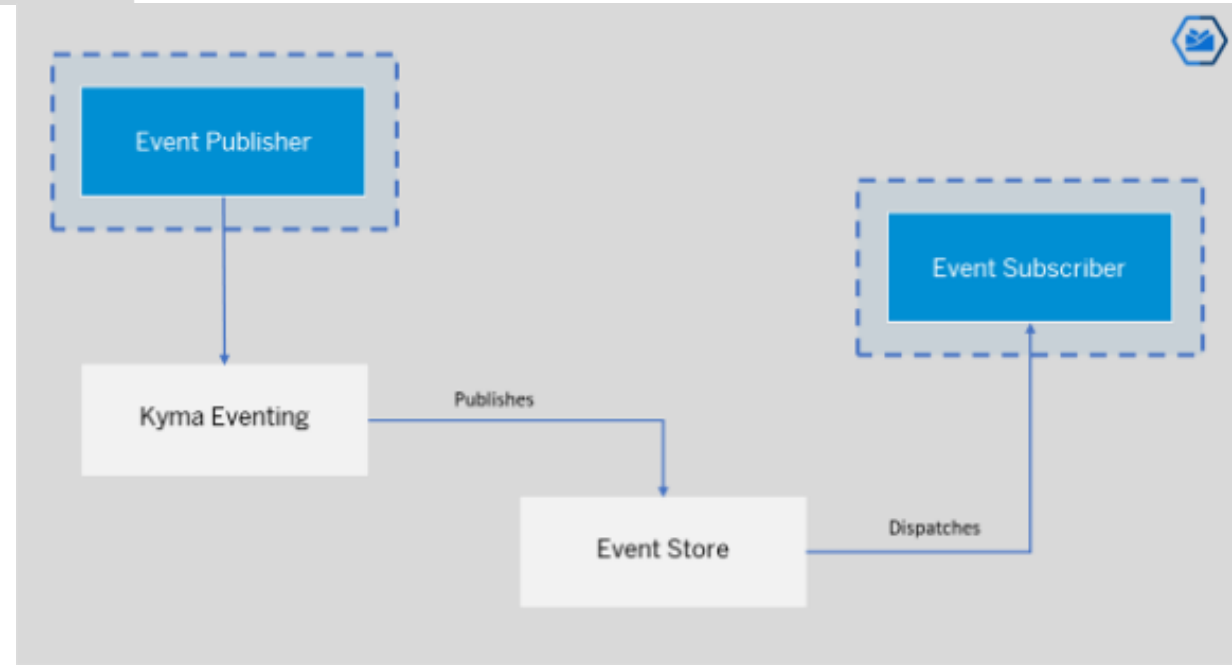


Unit 6 – Kyma Eventing

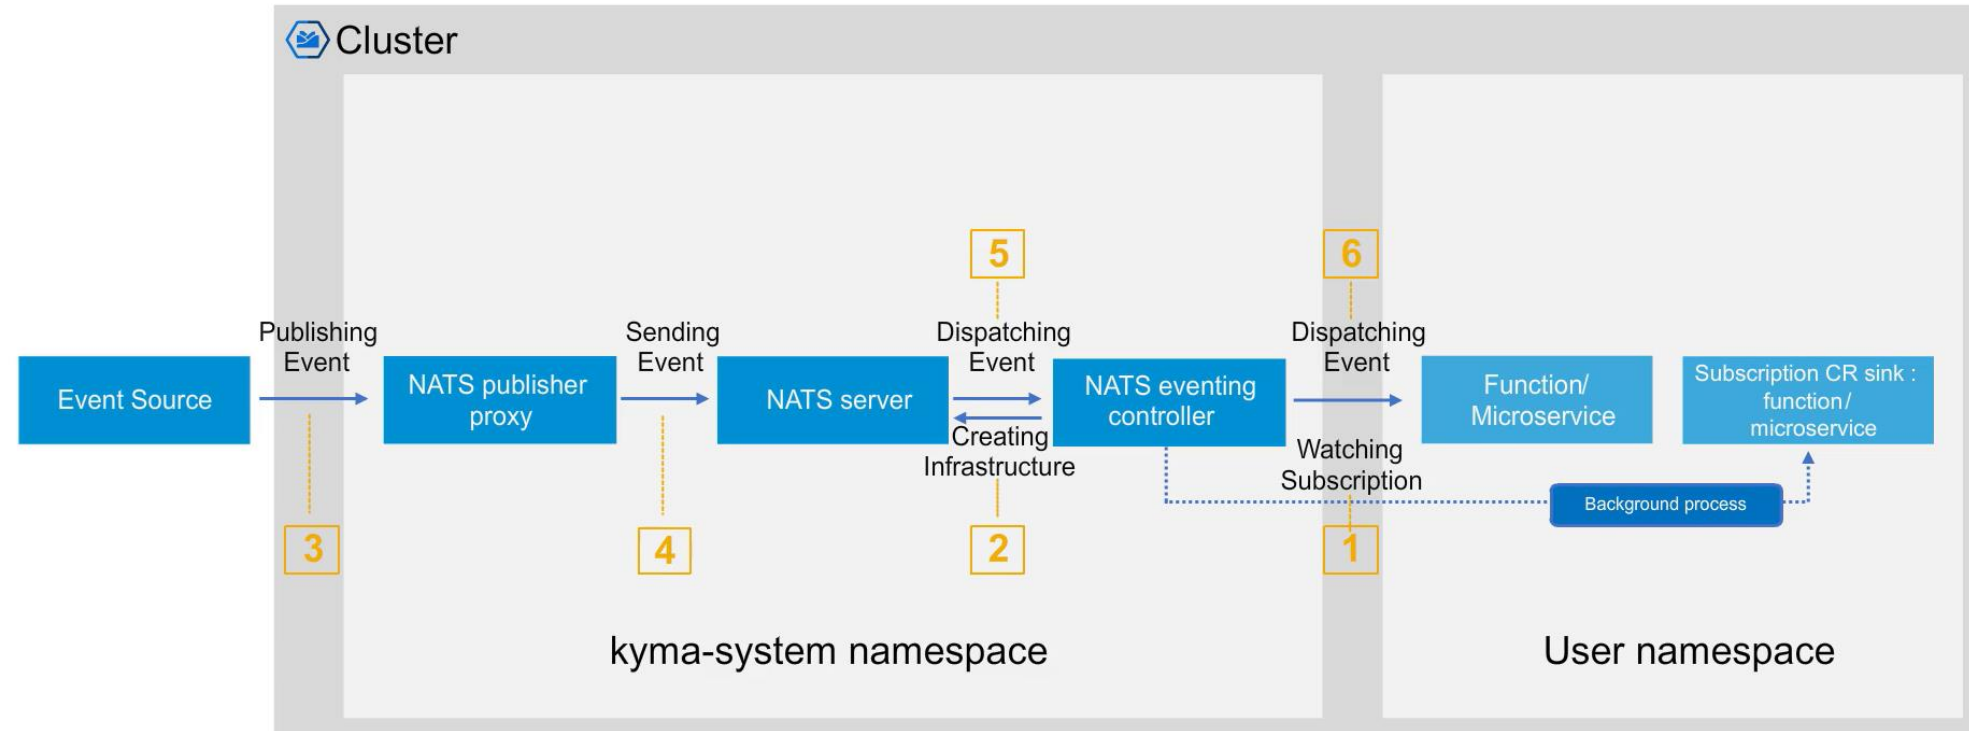


Business events are consumed from connected SAP and non-SAP solutions

In-Cluster events are used for event driven microservices



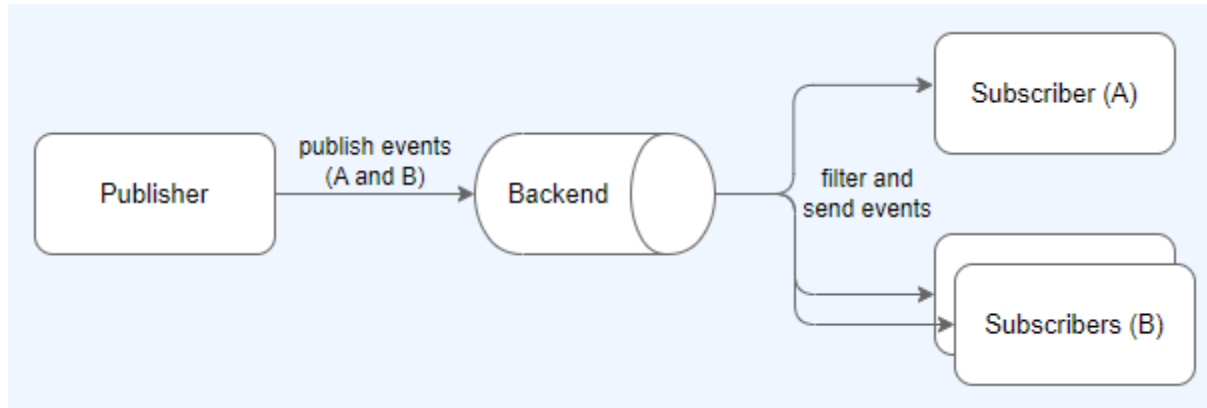
Unit 6 – Kyma Eventing



Kyma Eventing process uses

- NATS (an open source messaging system)
- NATS JetStream (provides the backend)
- HTTP post requests (for sending and receiving events)
- Subscription

Unit 6 – Eventing in Kyma



Eventing in Kyma – how it works ?

- Offer an HTTP endpoint, for example a function to receive events
- Specify the events the user is interested in using Kyma [Subscription CR](#)
- Send [CloudEvents](#) to the following HTTP endpoints on our Eventing Publisher Proxy Service
 - /publish for [CloudEvents](#)

Unit 7 – StatefulSet

In most cases, pods are stateless

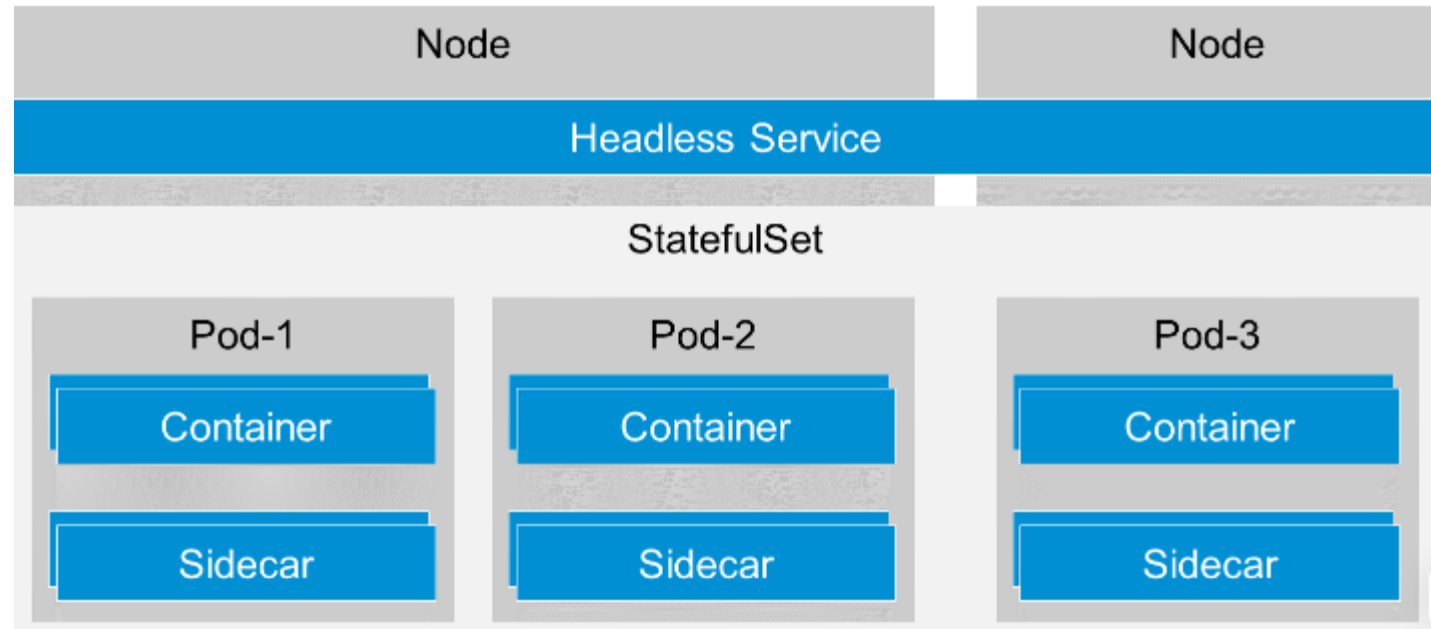
They die and are replaced **by any other pod with some random name**

But stateful applications (Database) need a unique, consistent identifier.

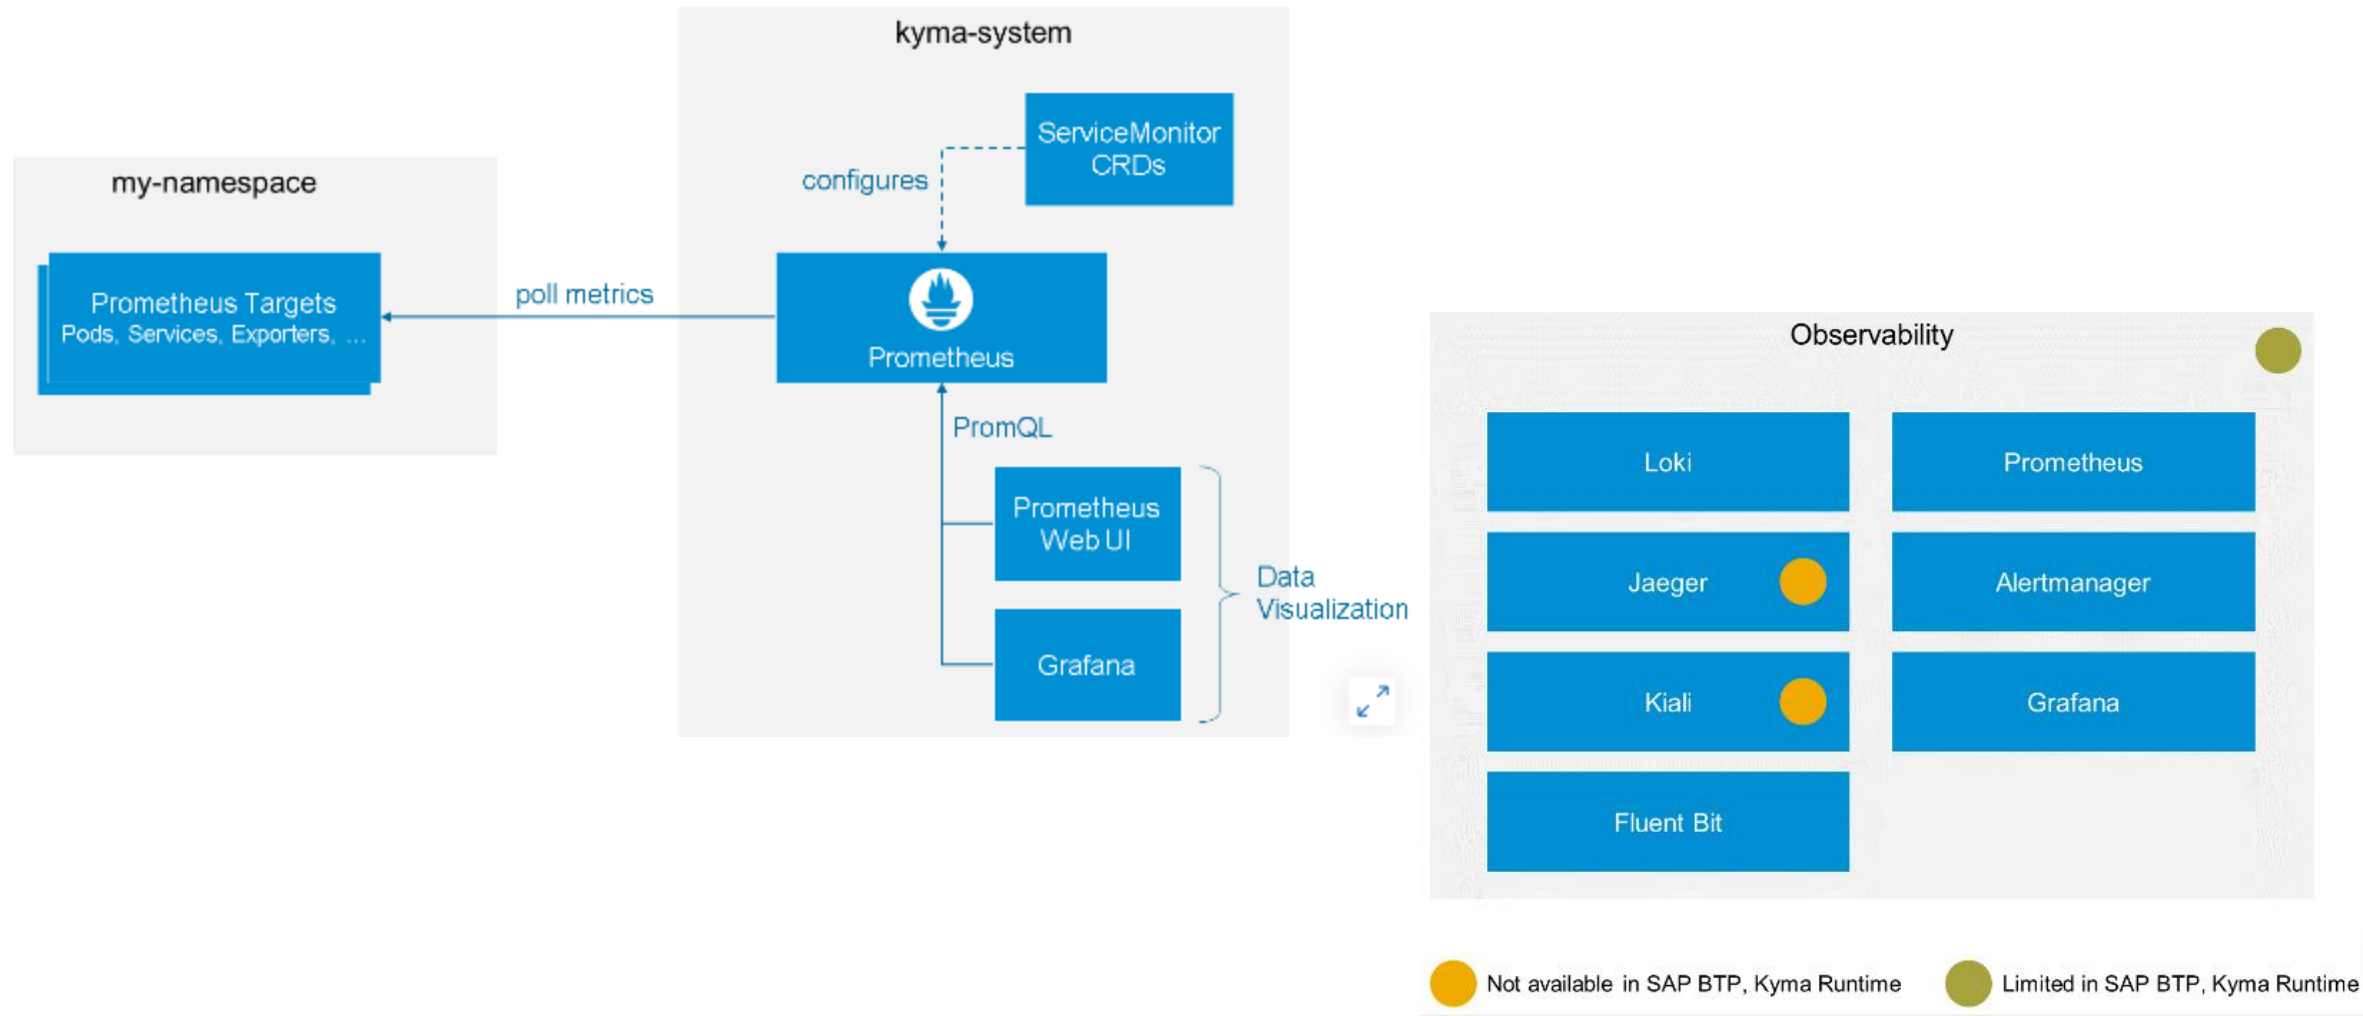
Use StatefulSets when you have the following requirements (e.g. Database)

- Stable, unique network identifiers
- Stable, persistent storage
- Ordered, graceful deployment and scaling
- Ordered automated rolling updates

When creating a StatefulSet, pods are created in sequential order (not any random order)

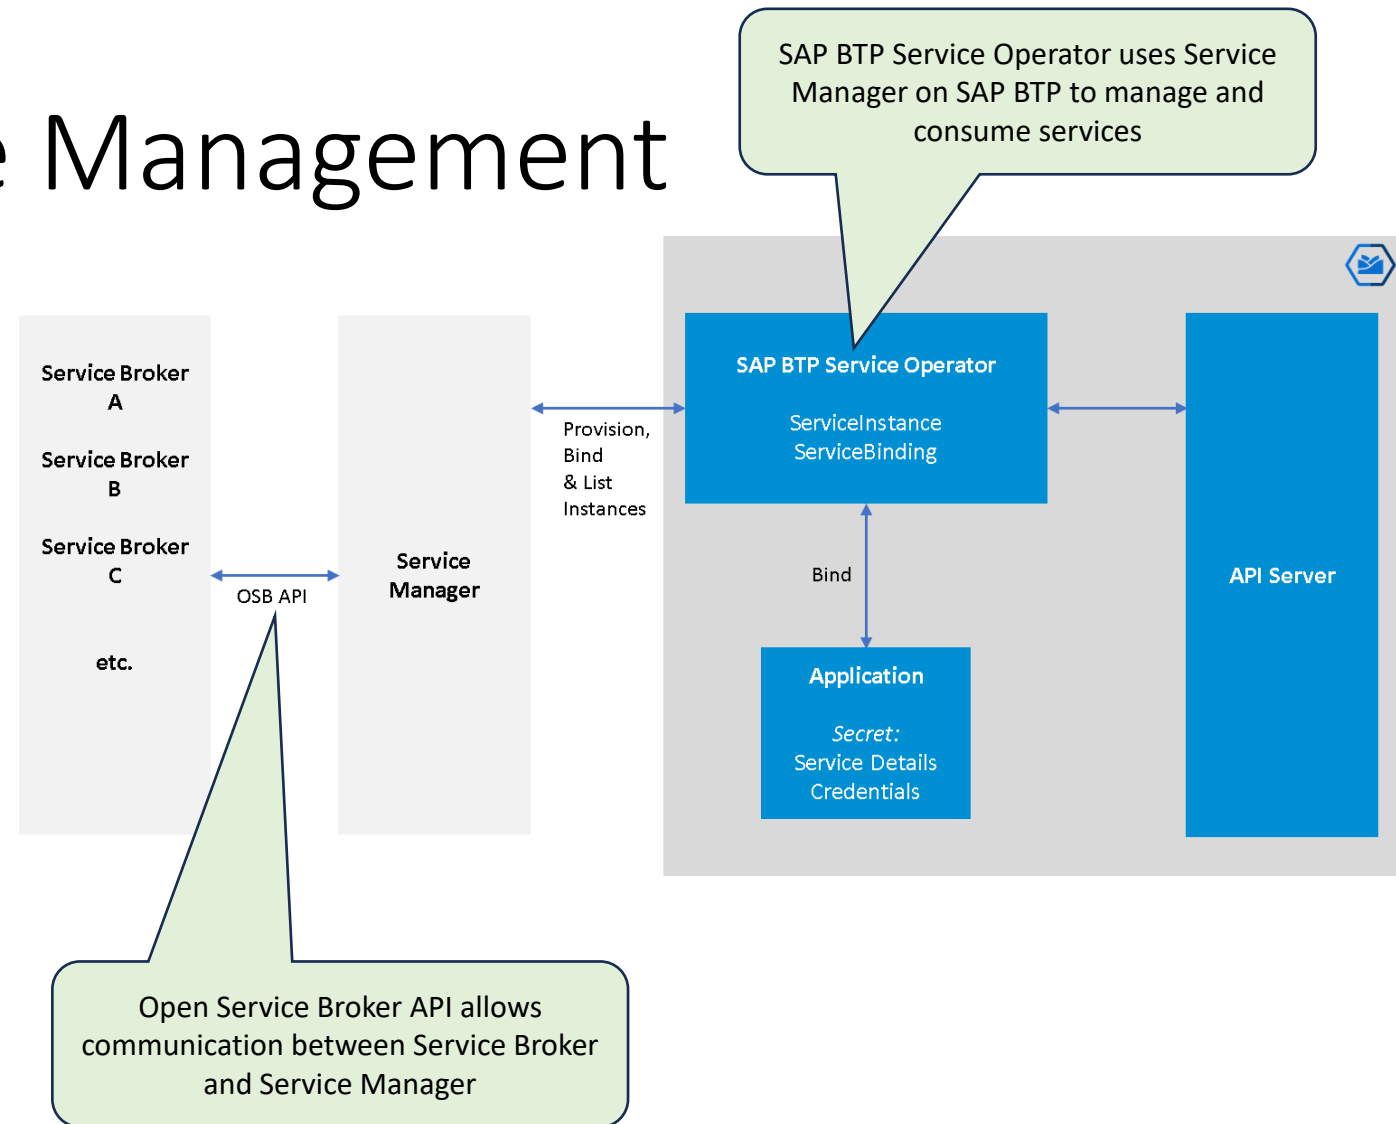


Unit 8 – Observability and Monitoring



Unit 9 – BTP Service Management

Service Management allows you to connect to SAP BTP services to your cluster



Unit 9 – SAP BTP Service Management



Unit 10 – ConfigMaps and Secrets

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: name-greeting
5  data:
6    GREETING: "Hello"
7    FIRSTNAME: "Kyma"
```



K8s has 2 native resources to store configuration data decoupled from application

ConfigMaps

Non Sensitive data

Secrets

Sensitive data (passwords, apikey etc.)

Unit 10 – Config Maps, Secrets

